

Job Preferences, Labor Market Power, and Inequality

Oscar Volpe

November 12, 2025

[Access the [latest draft](#) and [online appendix](#).]

Abstract

I analyze how a firm's labor market power shapes, and is shaped by, its workforce, and I evaluate the implications for welfare and inequality. Using matched worker-firm panel data from Norway (1995-2018), I develop, identify, and estimate an equilibrium model of the labor market where firms compete with one another for workers who are heterogeneous in both their skills and preferences over wages versus non-wage job amenities. I allow the wage-amenity trade-offs to be correlated with skills, while also varying among equally skilled workers. When a firm adjusts its wages, the composition of its workforce shifts, and these compositional changes, in turn, affect the labor supply curve to the firm. As a result, the firm's wage-setting power varies based on which types of workers it employs. I find that this variation leads to large allocative inefficiency, with welfare losses from imperfect competition estimated at 9.5% relative to the competitive benchmark.

I. Introduction

The role of firms in shaping wage inequality is a central topic for labor policy. Three growing empirical literatures examine how the characteristics of firms influence workers' earnings. One body of research (Manning, 2021; Card, 2022; Azar & Marinescu, 2024) studies the role of imperfect competition in the labor market, where firms exert wage-setting power due to facing upward-sloping labor supply curves. A second literature—building on work by Abowd, Kramarz, & Margolis (1999)—focuses on sorting, using matched employer-employee data to uncover the relationship between worker and firm productivity. A third literature examines

I am grateful to my advisors Alex Torgovitsky, Magne Mogstad, Steven Durlauf, and Thibaut Lamadon for their invaluable guidance and encouragement. I also thank Alan Manning, Evan Rose, Derek Neal, Manasi Deshpande, Michael Dinerstein, Michael Greenstone, Mikhail Golosov, Stéphane Bonhomme for detailed and insightful comments, as well as Arnstein Vestre, Jeanne Sorin, Hugo Lopez, Lancelot H. de Frahan, Michael Gruber, Ola Vestad, Tanya Rajan, Tom Melling, and the participants at the Applied Micro working groups and Public/Labor lunches at the University of Chicago for helpful feedback and suggestions. All errors are my own.

*Department of Economics, University of Chicago.

workers' willingness to trade higher wages for better non-wage aspects of a job, finding that the trade-offs vary substantially among workers (Mas & Pallais, 2017; Wiswall & Zafar, 2018).

While often studied separately, the issues explored by these literatures are interconnected. In a labor market with differentiated employers, a firm's wages and non-wage amenities determine which workers it attracts. Furthermore, if workers make different trade-offs between wages and amenities, then the composition of a firm's workforce could matter for the elasticity of labor supply that it faces, which in turn impacts the firm's wage-setting power. These factors prompt a key question: *How is a firm's wage-setting power linked to its workforce?* Addressing this question would offer deeper insight into how workers sort and how market power varies within the labor market, which has important implications for welfare and inequality.

This paper develops an empirical framework to examine how workers choose jobs based on heterogeneous wage-amenity trade-offs, and the consequences of this behavior for a firm's wage-setting power. I show that the market power at a firm—as measured by wage markdowns and rents—directly depends on the composition of its workforce, which is influenced by the wages and amenities the firm provides. Moreover, by failing to account for this relationship, one would overlook a key source of allocative inefficiency that arises from differences in wage markdowns among heterogeneous firms. I show how to recover economic quantities under my framework using matched worker-firm panel data, while allowing for unobserved heterogeneity in worker skills, firm productivity, technology, and amenities. I then apply my method to administrative data from Norway, where I draw inference about the determinants of worker sorting, the concentration of market power, and the welfare impacts of imperfect competition.

I begin my analysis by presenting an equilibrium model of the labor market. Building on previous models of job differentiation, such as Card et al. (2018), Lamadon et al. (2022), and Azar et al. (2022), I assume that workers' idiosyncratic tastes are not fully priced into wages, leading firms to face imperfectly elastic labor supply curves. However, unlike previous work, I allow workers to have varying marginal rates of substitution between wages and amenities. Drawing on the literature about the willingness to pay for amenities, I allow the wage-amenity trade-offs to be correlated with skills, while also differing among equally skilled workers.

This extension has key implications for the shape of the labor supply curve faced by a firm. It establishes a link between worker and firm characteristics, such that the elasticity of labor to the firm depends on the composition of its workforce. Consequently, the elasticities differ among firms, and they are also endogenous to changes in the environment: if a firm adjusts its wages, then the composition of its workforce would shift, and these compositional changes would, in turn, influence the elasticities. This variation contrasts with the standard assumption of homogenous and isoelastic labor supply curves, which leads to uniform wage markdowns. Moreover, to the extent that workers' wage-amenity trade-offs are correlated with their skills, wage markdowns may also differ within a firm—among workers with varying

skills—as well as across firms. These implications do not rely on firms competing strategically for workers, as in Berger et al. (2022). Even when firms are strategically small within the labor market, their wage-setting power would still vary based on *who* is employed at each firm.

To account for many potential sources of worker sorting, I allow firms to be differentiated along several key dimensions. Specifically, I allow firms to vary in productivity, such that a worker’s skills may be valued differently at different firms. I also allow for heterogeneous technology in the production function, where firms can exhibit decreasing returns to scale and imperfect substitutability between labor inputs. Moreover, I assume that firms have distinct non-wage amenities, which also differ within a firm based on workers’ skills.¹ I demonstrate how each of these firm characteristics interacts with workers’ preferences in equilibrium, and I discuss the consequences of these interactions for wage markdowns, rents, and misallocation.

I then turn to the question of identification, demonstrating how to learn about imperfect competition under my model using matched worker-firm data. In my analysis, identification is challenging for two reasons. First, as is typical in most contexts, a worker’s wage markdown is not observed in the data. In particular, observationally similar workers could earn different wages due to unobserved skill differences—not market power.² Wages may also not fully reflect a worker’s compensation, as firms may have amenities that researchers do not observe. Indeed, as Rosen (1986) discusses, compensating differentials can exist even in competitive markets. A second identification challenge stems from the generality of my framework, which allows workers’ wage markdowns to vary both within and across firms, and to be endogenous to wage changes within a firm. This added feature complicates my analysis because instead of recovering a single markdown, I need to recover entire functions of firm characteristics.

I show how to overcome these challenges to achieve identification of the model through the use of panel data by leveraging an economic shock that affects the productivity of multiple firms in the economy. These productivity shocks can be recovered either from existing data using internal instruments or from external data sources. One novelty of my approach is that I allow the productivity shocks to shift wages in a way that is reflected in workers’ wage indices. Therefore, I allow there to be spillover effects on the labor supply curves faced by *all* firms, including those that do not experience a productivity shock. These spillovers would prevent me from recovering the elasticity of labor supply to a firm by directly comparing a firm’s outcomes before and after the shocks. Nevertheless, I prove that—under weak conditions on

¹Recent empirical work provides mixed evidence about which non-wage job characteristics are most important to workers; for discussion, see Maestas et al. (2017). In my analysis, I do not take a stance on this issue. Instead, I assume that firms are endowed with a fixed set of amenities that are unobserved by the analyst; or, more precisely, I restrict amenities to be fixed over the estimation window. This assumption neither imposes nor precludes the possibility that firms initially choose amenities to maximize profits. Indeed, it is straightforward to show that permitting firms to choose their amenities initially would have no impact on any of my estimates.

²See, e.g., Murphy & Topel (1990), Gibbons & Katz (1992), Abowd et al. (1999), and Gibbons et al. (2005).

the instrument—the elasticities are identified from a difference-in-differences comparison of outcomes. This method involves comparing mean changes in wages and labor between firms with different exposures to productivity shocks, while controlling for a firm’s initial wages.

An additional contribution of my analysis is that I draw inference about the differences in workers’ trade-offs between wages and amenities. This approach is based on a revealed preference argument: holding all else fixed, a firm that offers better wages (for a given skill group) is more likely to attract workers who value wages more highly. This compositional effect is captured by the labor supply elasticity faced by a firm. Given this property, I prove that there exists a one-to-one mapping between the density function governing workers’ wage-amenity trade-offs and a firm’s labor supply curve. I then develop a method to estimate this density nonparametrically, which can be implemented separately for different population subgroups.

Using this identification strategy, I demonstrate how to recover the structural parameters in the model that characterize firm productivity, technology, and amenities, as well as workers’ preferences. From these parameters, I can calculate the rents for both workers and firms, and conduct counterfactual analyses to explore sources of allocative inefficiency and inequality.

Finally, I estimate the model using a matched employer-employee panel dataset, covering the universe of workers and firms in Norway for the period 1995-2018. My results yield three key findings. First, I find that many firms have substantial wage-setting power. On average, workers’ wages constitute only 86% of their marginal products. Workers and firms also earn large rents: on average, workers are willing to forgo 17% of their wages to remain at their current job, while the average firm derives 29% of its profits by exerting wage-setting power.³

Second, I find that labor market power varies significantly within and across firms, with the wage-to-marginal-product ratio ranging from 0.70 to 0.95. I also find that market power is more concentrated among firms that pay lower wages for a given skill group. This finding contrasts with the traditional view that larger, more productive firms should have more wage-setting power over workers.⁴ However, it supports a key prediction of my model. Specifically, it suggests that lower-paying firms are more likely to employ workers who are less responsive to wage changes (less elastic), thereby enabling these firms to exert greater wage-setting power.

Third, I find that imperfect competition generates substantial misallocation of workers to firms. By eliminating labor wedges, I estimate that total welfare would increase by 9.5%, measured in dollar terms. Moreover, in a competitive (Walrasian) labor market, more workers

³Empirically, it is hard to credibly differentiate between ex ante and ex post firm rents. This task necessitates further information—or assumptions—about how firms choose and pay for the amenities provided to workers.

⁴This view traces back to the early research on monopsony (Robinson, 1933), which examines a labor market with a single buyer of labor. However, if labor markets contain many firms, then market power need not increase with a firm’s size (Manning, 2021). Indeed, my estimates reveal that less productive firms tend to have greater wage-setting power because they are more likely to attract workers who are less responsive to wage changes.

would be employed at firms with lower productivity. This key source of inefficiency would not be captured by frameworks that do not account for heterogeneity in workers' wage markdowns.

My paper contributes to three literatures. First, it relates to a large literature about monopsony in the labor market; see Manning (2021), Sokolova & Sorensen (2021), Card (2022), and Azar & Marinescu (2024) for recent surveys.⁵ Within this literature, my paper builds on a class of static equilibrium models studied by Card et al. (2018), Azar et al. (2022), Lamadon et al. (2022), Roussille & Scuderi (2023), Chan et al. (2024), and Kroft et al. (2025), where imperfect competition arises from workers' idiosyncratic tastes for jobs not being fully priced into earnings. This work, which Manning (2021) terms "new classical monopsony," draws on differentiated demand models commonly employed in Industrial Organization. A limitation of these frameworks, however, is that they specify utility functions that strongly restrict how worker and firm characteristics interact. These restrictions constrain worker sorting patterns in a way that limits the potential for firm wage-setting power to vary in the economy.⁶ I relax these restrictions to allow for fully unrestricted and unobserved heterogeneity in workers' wage-amenity trade-offs. This extension is similar to one that is already well-established in the context of product markets, based on work developed by Berry et al. (1995). Moreover, as I demonstrate, it substantially expands the scope for worker sorting in labor markets. In doing so, it leads to a richer characterization of the determinants of firm wage-setting power.

Second, my paper contributes to an empirical literature on sorting, drawing on the additive worker and firm effects wage model introduced by Abowd, Kramarz, & Margolis (1999). This work has largely focused on skill-based sorting, arising, e.g., from worker-firm production complementarities (Song et al., 2018). While my model allows for sorting on workers' skills, I also study a second type of sorting based on workers' wage-amenity trade-offs. I show that this feature introduces a firm-skill interaction term—the wage markdown—in a worker's wage equation, such that the same firm can have a larger markdown for one skill group and a smaller markdown for another. In doing so, my paper provides additional insight into assortativeness in the allocation of workers to firms, e.g., Andrews et al. (2008) and Card et al. (2013).

Third, my paper contributes to a growing literature on the willingness to pay for non-wage amenities. Both experimental studies and survey data offer strong evidence that workers value a wide variety of non-wage job attributes (Hamermesh, 1999; Pierce, 2001; Maestas et al., 2017). Additionally, the willingness to pay for these amenities differs among workers (Mas

⁵For additional discussion, I refer to Boal & Ransom (1997), Bhaskar et al. (2002), and Manning (2011).

⁶Abstracting from strategic interactions, it leads to uniform wage markdowns. As discussed in Section II.E, this restriction can be relaxed by using a nested logit model (Lamadon et al., 2022; Azar et al., 2022) or by letting markdowns vary within a firm based on observable worker characteristics (Chan et al., 2024; Deb et al., 2024). In testing firm conduct assumptions, Roussille & Scuderi (2023) also allow for worker-firm interactions, but limit worker heterogeneity by assigning workers to a finite set of types, with selection into each type determined by a constant linear function of observables plus a logistic error term. These approaches fail to capture sorting patterns that can arise from fully unobserved heterogeneity in both workers' wage-amenity trade-offs and skills.

& Pallais, 2017; Wiswall & Zafar, 2018), contributing to wage inequality (Maestas et al., 2023). My paper takes this research area in a new direction, exploring how differences in the valuations of amenities affect the wage-setting power of firms in imperfectly competitive labor markets. I demonstrate that workers' wage-amenity trade-offs are reflected in the labor supply elasticities faced by firms. Moreover, by recovering these elasticities, I can draw inference about differences in wage-amenity trade-offs among workers. This method provides a new way to learn about variation in the willingness to pay for amenities using revealed preference, contributing to a recent literature in this area, e.g., Sorkin (2018) and Hall & Mueller (2018).

This paper is organized as follows. Section II presents the model and explores equilibrium properties related to sorting, imperfect competition, and welfare. Section III discusses the data sources and sample selection. Section IV outlines the identification strategy. Section V describes the estimation procedure, parameter estimates, and fit. Section VI presents the empirical findings about rents, misallocation, and inequality. Lastly, Section VII concludes.

II. Equilibrium Model of the Labor Market

In this section, I develop a theoretical framework to characterize how workers sort into jobs and the implications of these sorting patterns for a firm's wage-setting power. This framework extends prior equilibrium models of the labor market, studied by Lamadon et al. (2022), Azar et al. (2022), Kroft et al. (2025) and others, by allowing workers to differ in the marginal rate of substitution between wages and non-wage aspects of jobs. I show that this extension leads firms to face heterogeneous labor supply elasticities, which reflect differences in the composition of workers at different firms. As a result, this model generates variation in the incidence of wage markdowns and rents, thus introducing a potential source of allocative inefficiency.

II.A. Environment

The economy comprises a unit measure of workers, indexed by i , and a finite number of firms, indexed by j . Every firm operates in a local labor market $m(j)$, where \mathcal{J}_m is the set of firms in market m . Workers are differentially productive, and they have heterogeneous preferences over employers. Meanwhile, firms provide differentiated work environments, and they vary in their productivities and technologies. Note that, while distinguishing between labor markets will be important for the empirical analysis, it does not meaningfully impact the theoretical results. I therefore simplify notation by focusing on a single market and suppressing the label m .

Each worker has a vector of skills $X_i = (\chi_i, \varphi_i)$. These skills contain two components: a categorical *skill type* $\chi_i \in \mathcal{X}$ and a continuous *skill level* $\varphi_i \in \mathbb{R}_+$. The skill type χ_i captures non-ordinal aspects of human capital, such as training and specialization. The skill level φ_i , which I treat as unobservable to the researcher in the empirical analysis, measures productivity within a given skill type. By modeling skills as a composite of types and levels, the framework

incorporates insights from the literature on multidimensional skills (e.g., Lise & Postel-Vinay, 2020), while also allowing human capital to vary in ways that a researcher cannot observe.

Throughout the analysis, I will maintain two assumptions about the information structure. First, I abstract from search frictions by assuming that workers are fully informed about job opportunities and that, given posted wages, a worker freely chooses where to work. Second, I assume that firms can observe workers' skills, but they only know the distribution of workers' preferences conditional on those skills. Therefore, while firms can set wages based on X_i , they cannot further price-discriminate with respect to each worker's idiosyncratic preferences.⁷

II.B. Worker Preferences

Workers have heterogeneous preferences over wages and non-wage amenities. Let $W_j(X)$ be the wage that a firm j offers to workers with skills X , and let $a_j(X)$ be a firm's skill-specific amenity. For any worker i with skills X_i , the indirect utility from working at firm j is:

$$u_{ij}(W_j(X_i), a_j(X_i)) = \beta_i \log W_j(X_i) + a_j(X_i) + \epsilon_{ij}. \quad (1)$$

In this utility function, ϵ_{ij} is a worker's idiosyncratic taste for firm j , reflecting non-pecuniary match effects such as distance from work and personal preferences over work environment. I assume $\{\epsilon_{ij}\}_{i,j}$ are i.i.d. with a continuous distribution F_ϵ taking strictly positive density on \mathbb{R} .

The coefficient β_i is the worker's marginal utility of log earnings. This parameter governs the marginal rate of substitution between wages and amenities, thereby capturing the trade-off between higher pay and better working conditions. Workers with lower β_i will be more willing to sacrifice earnings in exchange for better non-wage job characteristics. In Appendix A.9, I provide an explicit micro-foundation for the utility function (1), where I show that β_i can be reinterpreted as the worker's marginal rate of substitution between consumption and leisure.

In my analysis, I allow β_i to vary freely among workers. In particular, I treat β as a random variable with density f_β , and I do not impose any parametric restrictions on this density, such as assuming normality. One important aspect of this framework is that I allow for two distinct forms of heterogeneity in β . First, I allow β to correlate with a worker's skills X , so that the trade-offs workers make when choosing jobs can depend on factors that impact productivity. Second, I allow β to vary conditionally within each skill group X , so that equally skilled workers can have different wage-amenity trade-offs. This within-skill heterogeneity is particularly important, as it allows worker's trade-offs to reflect considerations beyond productivity. For example, β may depend on factors such as age, gender, household characteristics, and wealth. It may also be correlated among workers in the same household, neighborhood, or peer group.

⁷This assumption reflects information asymmetries between workers and firms. It may also account for laws against pay discrimination, which are common in many labor markets. See Bhaskar et al. (2002) for discussion.

II.C. Firm Production

Let $D_j(\chi, \varphi)$ denote the mass of workers with skills $X = (\chi, \varphi)$ that is demanded by firm j . For each skill type $\chi \in \mathcal{X}$, the total efficiency units of labor at the firm are defined as follows:

$$L_j^{\text{eff}}(\chi) = \int \varphi D_j(\chi, \varphi) d\varphi. \quad (2)$$

Define Y_j to be the firm's value added from hiring workers. The firm's production function is a constant elasticity of substitution (CES) composite of the effective labor from different skill types. This specification permits imperfect substitutability between skill types and allows for decreasing returns to scale, so that output need not increase in direct proportion to labor inputs.

$$Y_j = T_j \left(\sum_{\chi \in \mathcal{X}} \theta_{j\chi} \left(L_j^{\text{eff}}(\chi) \right)^{\rho_j} \right)^{\frac{1-\alpha_j}{\rho_j}}. \quad (3)$$

In this production function, $T_j \in \mathbb{R}_{++}$ is the firm's total factor productivity (TFP); $\theta_{j\chi} \in \mathbb{R}_{++}$ is the firm-specific efficiency of skill type χ ; $(1 - \rho_j)^{-1} \in [0, \infty)$ is the elasticity of substitution between skill types at the firm; and $1 - \alpha_j \in (0, 1]$ is the firm's returns to scale. For simplicity, I abstract from capital and intermediate inputs, as well as from the nature of competition in the product market. Nevertheless, as I show in Appendix A, this production function can be rationalized under either a perfectly competitive or an imperfectly competitive product market.

By allowing the parameters $(T_j, \{\theta_{j\chi}\}_{\chi \in \mathcal{X}}, \rho_j, \alpha_j)$ to vary among firms, the model accommodates a very wide range of production technologies. First, firms can differ in their overall productivity through T_j . Second, firms can differ in the relative efficiency of each skill type through $\{\theta_{j\chi}\}_{\chi \in \mathcal{X}}$; for example, some firms might place more value on a college degree than others. Third, firms can differ in the degree of substitutability between skill types through ρ_j ; for example, some firms might have specific job requirements that affect their ability to interchange high- and low-skill workers. Fourth, firms can exhibit different returns to scale through α_j . Importantly, I do not place any restrictions on the relationship between amenities, productivity, and technology. Rather, I allow these fundamentals to vary freely across firms.

II.D. Characterization of an Equilibrium

In equilibrium, each worker chooses a firm to maximize utility. Given wage offers $\{W_k(X)\}_{k=1}^J$, a worker with preference parameter β and skill profile X would select firm j with probability:

$$P(j(i) = j | \beta, X) = \int_{-\infty}^{\infty} \prod_{k \neq j} F_\epsilon \left(\beta \log \left(\frac{W_j(X)}{W_k(X)} \right) + a_j(X) - a_k(X) + \tilde{\epsilon} \right) f_\epsilon(\tilde{\epsilon}) d\tilde{\epsilon}. \quad (4)$$

Averaging over realizations of β , the labor that is supplied to firm j by workers with skills X is:

$$S_j(X) = \int \left(\int_{-\infty}^{\infty} \prod_{k \neq j} F_{\epsilon} \left(\beta \log \left(\frac{W_j(X)}{W_k(X)} \right) + a_j(X) - a_k(X) + \tilde{\epsilon} \right) f_{\epsilon}(\tilde{\epsilon}) d\tilde{\epsilon} \right) f_{\beta,X}(\beta, X) d\beta. \quad (5)$$

Given these labor supply curves, each firm sets wages to maximize profit. Since a firm cannot condition wages on individual preferences (β_i, ϵ_{ij}) , its wage-setting decisions are based on the distributions of these parameters, $\{F_{\beta|X}\}_X$ and F_{ϵ} . The firm's profit maximization problem is:

$$\max_{\{W_j(\chi, \varphi)\}_{\chi, \varphi}} T_j \left(\sum_{\chi \in \mathcal{X}} \theta_{j\chi} \left(\int \varphi D_j(\chi, \varphi) d\varphi \right)^{\rho_j} \right)^{\frac{1-\alpha_j}{\rho_j}} - \sum_{\chi \in \mathcal{X}} \left(\int W_j(\chi, \varphi) D_j(\chi, \varphi) d\varphi \right), \quad (6)$$

subject to the constraint that labor demand equals supply: $D_j(X) = S_j(X)$ for all $X = (\chi, \varphi)$.

I consider an equilibrium in which firms view themselves as strategically small in the economy. Specifically, the firm does not internalize the effect of changing its own wages on the labor supplied to other firms. It is worth noting that this condition does not imply that workers regard firms as infinitesimal. On the contrary, there is a finite number of firms, and each worker chooses a firm by considering the entire set of offered wages and amenities. Instead, this condition applies only to firm conduct; it assumes that a firm sees itself as a marginal player, thus ruling out strategic interactions in wage-setting. This assumption also appears in Card et al. (2018), Lamadon et al. (2022), Azar et al. (2022), and Kroft et al. (2025), among others.⁸ In Section V, I derive a statistical test of this assumption. Applying this test to Norwegian administrative data, I find no evidence to reject that firms view themselves as strategically small.

Definition. Given worker skill and preference distributions $(F_{\chi, \varphi}, \{F_{\beta|\chi, \varphi}\}_{\chi, \varphi}, F_{\epsilon})$, firm amenities $\{a_j(\chi, \varphi)\}_{\chi, \varphi}$ and production parameters $(T_j, \{\theta_{j\chi}\}_{\chi}, \rho_j, \alpha_j)$, an equilibrium is defined by the employment outcomes $j(i)$, labor supply curves $S_j(\chi, \varphi)$, and wages $W_j(\chi, \varphi)$ where:

- (i) Each worker chooses the firm that maximizes utility, with utility given by equation (1).
- (ii) Labor supply curves $S_j(\chi, \varphi)$ are consistent with workers' optimal decisions, as in (5).
- (iii) Each firm sets wages $W_j(\chi, \varphi)$ for all $\chi \in \mathcal{X}$ and $\varphi \in \mathbb{R}_+$ to maximize profit, as defined in equation (6), ensuring that labor demand $D_j(\chi, \varphi)$ is equal to labor supply $S_j(\chi, \varphi)$.

In Lemma 1 of Appendix B.1, I establish the existence of an equilibrium with strictly positive wages and employment. Lemmas 2 and 3 further characterize properties that relate to equilibrium uniqueness. Additionally, I analyze properties such as the concavity of the profit function

⁸ Alternatively, Berger et al. (2022), Deb et al. (2024), and Chan et al. (2024) study models of job differentiation where firms internalize how their wages affect market aggregates, and Jarosch et al. (2024) analyzes a search model with large firms. Roussille & Scuderi (2023) develop a procedure to compare models of firm conduct and find that, in their context, models which ignore strategic interactions outperform those that incorporate them.

and dynamic stability of the equilibrium. These results are useful for counterfactual analysis, as they enable me to compute equilibrium outcomes from the underlying model parameters.

II.E. Properties of the Labor Supply Curve to a Firm

A primary motivation for incorporating random coefficients in the worker's utility function is that it enables the model to capture rich substitution patterns that reflect sorting on individual heterogeneity. This subsection examines how this heterogeneity expands the scope of worker behavior in the model and what it implies about the shape of the labor supply curve to a firm.

II.E.1. Labor Supply under Homogeneous Effects

The standard assumption in models of job differentiation is that β is constant across workers. Under this assumption, firm wages and amenities do not interact with worker-specific preferences in the utility function. This additive separability between firm and worker characteristics delivers a relatively simple and tractable labor supply system. Yet, it imposes a strong form of homogeneity on aggregate substitution patterns, which strongly restricts worker sorting under the model. This critique has been raised in the context of product markets (Berry et al., 1995). Moreover, as I will now demonstrate, it also holds significant implications for labor markets.

To illustrate how the homogeneous effects assumption restricts economic behavior, consider a version of the model where β is constant across workers. In this case, workers differ in how they rank firms only through the idiosyncratic taste shocks ϵ_{ij} . Because these shocks are i.i.d., a worker's second-choice firm is independent of her chosen firm. This restriction has strong implications for how labor supply responds to a change in the economic environment. In particular, if wages shift in the market, a firm's employment response would not depend on its own characteristics. For example, consider two firms j and j' that employ the same share of workers with skill type X , such that $S_j(X) = S_{j'}(X)$. Let $S_{\mathcal{J}^*}(X) = \sum_{k \in \mathcal{J}^*} S_k(X)$ denote the labor supplied to any third set of firms \mathcal{J}^* , where $j, j' \notin \mathcal{J}^*$. If β is homogeneous, then:

$$\frac{\partial \log S_{\mathcal{J}^*}(X)}{\partial \log W_j(X)} = \frac{\partial \log S_{\mathcal{J}^*}(X)}{\partial \log W_{j'}(X)}, \quad \text{for any } \mathcal{J}^* \in \{1, \dots, J\} \text{ where } j, j' \notin \mathcal{J}^*. \quad (7)$$

This restriction implies that the workers at firms j and j' would be just as likely to substitute toward any third firm k , regardless of which firm is more similar to k . Indeed, even if j and j' have very different technologies and work environments, their employment responses to a wage change would be identical. This contradicts the standard intuition that firms with similar attributes tend to attract similar workers and should therefore exhibit larger substitution effects.

One way to relax this restriction on substitution patterns is to allow $\{\epsilon_{ij}\}_{i,j}$ to be correlated across firms in a local labor market, for example, via a nested logit structure. This approach is adopted by Azar et al. (2022), Lamadon et al. (2022), and others, and it implies that workers

are more likely to substitute toward firms within markets than across markets. Nevertheless, within any given labor market, the same restrictions on substitution patterns would still apply.⁹

A second extension could be to allow β to vary based on observable worker characteristics such as education or experience.¹⁰ Yet, a concern with this strategy is that firms are also likely to observe these characteristics and would thus factor them into a worker's skills X . As long as β is constant conditional on skills, the restrictions on substitution patterns remain unchanged.

II.E.2. Labor Supply under Heterogeneous Effects

If β is heterogeneous, then a firm's labor supply curve reflects differences in the composition of workers at different firms. For example, a firm offering higher wages and worse amenities is more likely to attract workers who place greater value on wages (higher β), whereas a firm offering lower wages and better amenities tends to attract workers who place greater value on amenities (lower β). Thus, firms with similar wage-amenity profiles are more likely to attract workers who make similar trade-offs, which leads these firms have larger substitution effects.

One key economic quantity that is affected by heterogeneity in β is the elasticity of labor supply to a firm, $\frac{\partial \log S_j(X)}{\partial \log W_j(X)}$. As shown in Lemma 4 of Appendix B.1, this elasticity can be represented as the average of β -specific elasticities among workers who are employed at the firm:

$$\frac{\partial \log S_j(X)}{\partial \log W_j(X)} = E\left(\frac{\partial \log P(j(i) = j | \beta_i, X_i)}{\partial \log W_j(X_i)} \middle| X_i = X, j(i) = j\right). \quad (8)$$

To understand how this labor supply elasticity varies across firms, I consider a version of the model where ϵ_{ij} follows a Type 1 Extreme Value (Logit) distribution, $F_\epsilon(\epsilon) = \exp(-\exp(-\epsilon))$. This is the most common parameterization of ϵ_{ij} in the literature, and it yields transparent expressions for the firm's labor supply curve, as well as its elasticities. In Lemma 5 of Appendix B.1, I show how the insights from the example below generalize to alternative error structures.

Example (Logit Model). In the logit model, a worker's choice probability, as defined in (4), is:

$$P(j(i) = j | \beta, X) = \frac{\exp(\beta \log W_j(X) + a_j(X))}{\sum_{k=1}^J \exp(\beta \log W_k(X) + a_k(X))}. \quad (9)$$

Define $I(\beta, X) = \sum_{k=1}^J \exp(\beta \log W_k(X) + a_k(X))$ to be the wage index for workers with preference parameter β and skills X . From the perspective of a strategically small firm, its wage has a negligible effect on $I(\beta, X)$. Thus, β coincides with the individual-specific labor supply elasticity, $\frac{\partial \log P(j(i)=j|\beta,X)}{\partial \log W_j(X)}$. The labor supply elasticity to a firm is then a firm-wide average of β :

⁹In principle, nests can be defined more broadly so they do not coincide with markets. However, such definitions still rely on factors observable to the analyst, which is often undesirable. See Nevo (2000) for discussion.

¹⁰Examples of this approach include Chan et al. (2024) and Deb et al. (2024). The former uses a multinomial logit discrete choice model and the latter uses a representative agent CES model. As shown in Verboven (1996), Proposition 2, both modeling approaches lead to demand functions that are isomorphic at the market level.

$$\frac{\partial \log S_j(X)}{\partial \log W_j(X)} = E(\beta_i | X_i = X, j(i) = j). \quad (10)$$

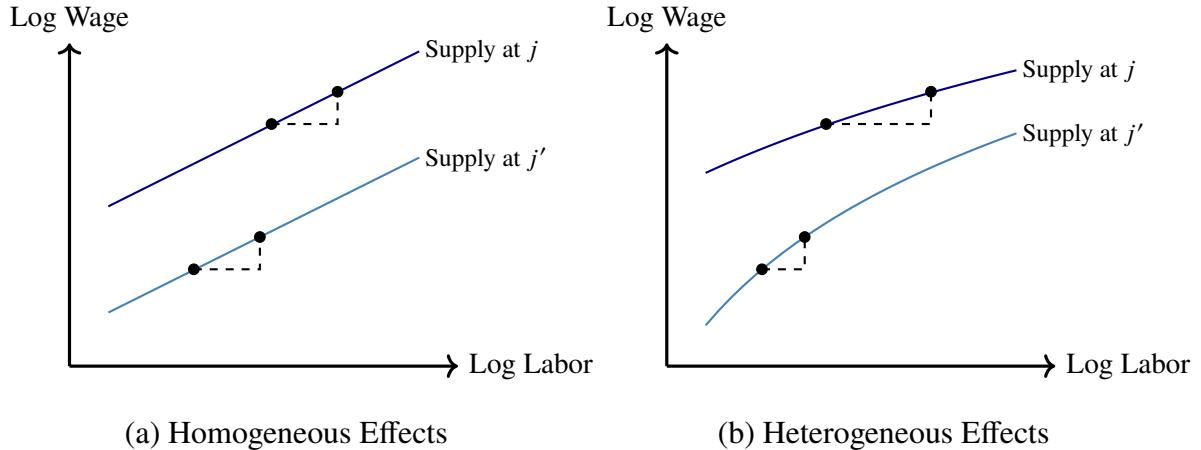
If β is homogeneous, this elasticity is the same for all firms. However, if β is heterogeneous, it varies across firms, because firms with different characteristics attract workers who make different trade-offs. As a result, each firm employs a workforce with a distinct distribution of β .

Another consequence of heterogeneous effects is that a firm's labor supply elasticity will be endogenous to its wage. When a firm posts a higher wage, it attracts workers who are more responsive to wages (higher β). As a result, this firm faces a more elastic labor supply curve. To formalize this dynamic, I analyze the second derivative of the labor supply curve to a firm:

$$\frac{\partial^2 \log S_j(X)}{\partial [\log W_j(X)]^2} = \text{Var}(\beta_i | X_i = X, j(i) = j). \quad (11)$$

This second derivative is strictly positive if and only if $\text{Var}(\beta_i | X_i = X) > 0$. Moreover, this quantity is strictly increasing in $\text{Var}(\beta_i | X_i = X)$, because greater heterogeneity in β expands the scope for sorting, making firm workforce composition more responsive to wage changes.¹¹

Figure 1: Illustration of Firm-Specific Labor Supply Responses



Notes. This figure depicts the labor supply curves $\log S_j(X)$ and $\log S_{j'}(X)$ for two firms, j and j' , respectively. Figure 1.a shows the case where $\text{Var}(\beta_i | X_i = X) = 0$. Figure 1.b shows the case where $\text{Var}(\beta_i | X_i = X) > 0$.

Figure 1 illustrates how allowing for heterogeneous effects impacts the labor supply curve to a firm. Consider two firms, j and j' , where firm j offers higher wages but worse amenities, and firm j' offers lower wages but better amenities. Under homogeneous effects (Figure 1.a), both firms face the same employment responses to an internal wage change, meaning that an

¹¹In Property A.1 of Appendix A, I derive the third, fourth, and fifth derivatives of a firm's labor supply curve under the logit model. I show that these quantities depend on the higher-order moments of the distribution of β .

$X\%$ wage change leads to the same percent change in labor at each firm. Under heterogeneous effects (Figure 1.b), firms face different labor supply responses due to differences in workforce composition. For the same $X\%$ wage change, the high-wage, low-amenity firm (j) sees a larger employment response, while the low-wage, high-amenity firm (j') sees a smaller response.

II.F. Wage Determination and Sorting in Equilibrium

In equilibrium, firms face upward-sloping labor supply curves, allowing them to set wages below workers' marginal products. Let $L_j(X)$ be the equilibrium employment of skill group X at firm j , and let $\varepsilon_j(X) = \frac{\partial \log L_j(X)}{\partial \log W_j(X)}$ be the labor supply elasticity firm j faces. A worker's wage is:

$$W_j(\chi, \varphi) = \underbrace{\frac{\varepsilon_j(\chi, \varphi)}{1 + \varepsilon_j(\chi, \varphi)}}_{\text{Markdown}} \times \underbrace{\varphi T_j(1 - \alpha_j) \theta_{j\chi} \left(L_j^{\text{eff}}(\chi)\right)^{\rho_j-1} \left(\sum_{\chi' \in X} \theta_{j\chi'} \left(L_j^{\text{eff}}(\chi')\right)^{\rho_j}\right)^{\frac{1-\alpha_j}{\rho_j}-1}}_{\text{Marginal Product of Labor}}. \quad (12)$$

There are three main factors that drive systematic worker sorting in the model. One factor is that workers' marginal products differ across firms. These differences arise from variation in firm productivity ($T_j, \{\theta_{j\chi}\}_\chi$) and technology (ρ_j, α_j). In particular, productivity differences imply that the same worker can be more productive at some firms than others, while technology differences affect how a worker's marginal product depends on the firm's existing labor inputs. Together, these forces have the potential to generate substantial dispersion in worker productivity within and across firms. Note that, with production complementarities ($\rho_j < 1$), wages need not be log-additively separable into worker and firm components, as is assumed in much of the sorting literature following Abowd, Kramarz, and Margolis (1999). By relaxing this log-additivity restriction, the model allows for more flexible sorting, since the same firm can generate a high marginal product for some workers and a low marginal product for others.

A second factor influencing sorting is that firms have distinct non-wage amenities $a_j(X)$, which also vary within a firm based on worker skills. These amenities create vertical firm differentiation, where some firms offer strictly better amenities than others. They also produce horizontal firm differentiation, where workers at the same firm benefit differently from the same amenity. Note that the amenities can depend arbitrarily on firm production parameters ($T_j, \{\theta_{j\chi}\}_\chi, \rho_j, \alpha_j$), which means they can be directly related to a worker's marginal product.

A third factor that affects sorting is the wage markdown $\frac{\varepsilon_j(X)}{1 + \varepsilon_j(X)}$. These markdowns differ across firms due to variation in workforce composition, reflecting an interaction between worker preferences and firm characteristics. For example, a low-productivity, high-amenity firm is more likely to attract workers who are less responsive to wages. This allows the firm to exert greater wage-setting power, leading to a larger wage markdown. In my model, the markdown is specific to each skill group X , since different skill groups can be differentially responsive to

wages. As a result, workers with different skills may face different markdowns even within the same firm. This again relaxes the log-additivity restriction of two-way fixed effects models, since a given firm may apply a high markdown to some workers and a low markdown to others.

II.G. Rents, Pass-throughs, and Allocative Inefficiency

I conclude this section by showing how the structural parameters in the model map into key economic quantities, such as rents earned by workers and firms, the pass-through of productivity shocks to earnings, and measures of welfare loss due to allocative inefficiency. To simplify these derivations, I assume a logit error structure for ϵ_{ij} , setting $F_\epsilon(\epsilon) = \exp(-\exp(-\epsilon))$. In each of the derivations in the Appendix, I carefully highlight which parts of the analysis rely on this assumption, and I also show how the results extend to alternative parameterizations of F_ϵ .

Worker Rents

In the economy, firms do not observe workers' individual preferences, β_i and ε_{ij} , which means they cannot price-discriminate based on each worker's reservation wage. Thus, an equilibrium involves surpluses—or rents—for inframarginal workers. These rents represent the earnings that workers are willing to forgo to be indifferent between their current firm and their next-best alternative. Let R_i^w denote a worker i 's rents at her current firm. I define these rents as follows:

$$u_{ij}(W_j(X_i) - R_i^w, a_j(X_i)) = \max_{j' \neq j(i)} u_{ij'}(W_{j'}(X_i), a_{j'}(X_i)). \quad (13)$$

In a model where β is homogeneous, a worker's rents as a share of earnings, $R_i^w/W_{j(i)}(X_i)$, are independent of her skills X and her firm j (see Property A.4 of Appendix A). Thus, a worker's employer and skills will have no impact on the fraction of wages she earns as rents. This places a strong restriction on the incidence of rents in the economy, and it runs counter to the idea that workers at different firms, or with different skills, benefit differently from their wage contracts.

When β is heterogeneous, the share of wages that workers earn as rents will systematically vary across firms and skill groups. To characterize this variation, I derive an expression for the average rent share, $E(R_i^w | j(i) = j, X_i = X)/W_j(X)$, among workers with skills X at firm j .

$$\frac{E(R_i^w | j(i) = j, X_i = X)}{W_j(X)} = \left(\frac{1}{1 + \beta_i} \middle| j(i) = j, X_i = X \right). \quad (14)$$

As shown in Property A.5 of Appendix A, this average rent share declines with the firm's wage:

$$\frac{\partial \left(\frac{E(R_i^w | j(i) = j, X_i = X)}{W_j(X)} \right)}{\partial \log W_j(X)} = \text{Cov} \left(\beta_i, \frac{1}{1 + \beta_i} \middle| j(i) = j, X_i = X \right) \leq 0. \quad (15)$$

This negative relationship reflects how workers sort based on individual heterogeneity. Holding skills fixed, higher-wage firms attract a larger share of workers who are more responsive

to wages. These workers would be more likely to switch firms in response to small pay differences. Therefore, even though they earn higher wages, a smaller share of those wages would constitute rents. As I demonstrate empirically in Section VI, neglecting this compositional effect leads to an inaccurate assessment of where worker rents are concentrated in the economy.

Firm Rents

Firms also capture rents in equilibrium. These rents come in the form of excess profit that firms obtain by exerting their wage-setting power. To quantify these rents, I analyze a counterfactual setting where a firm is a price-taker in the labor market, acting as if it faces a perfectly elastic labor supply curve. For each firm j , I construct these counterfactual profits Π_j^{pt} by solving:

$$\max_{\{D_j^{\text{pt}}(\chi, \varphi)\}_{\chi, \varphi}} T_j \left(\sum_{\chi \in \mathcal{X}} \theta_{j\chi} \left(\int \varphi D_j^{\text{pt}}(\chi, \varphi) d\varphi \right)^{\rho_j} \right)^{\frac{1-\alpha_j}{\rho_j}} - \sum_{\chi \in \mathcal{X}} \left(\int W_j^{\text{pt}}(\chi, \varphi) D_j^{\text{pt}}(\chi, \varphi) d\varphi \right), \quad (16)$$

subject to the labor supply constraint: $D_j^{\text{pt}}(X) = S_j^{\text{pt}}(X)$ for all $X = (\chi, \varphi)$. In this counterfactual, the only difference in the firm's behavior is that it does not exert its wage-setting power. Hence, this counterfactual is distinct from an equilibrium where all firms act as price takers.

I define a firm j 's rents as the difference between its realized profits and its counterfactual profits, given by $R_j^e = \Pi_j - \Pi_j^{\text{pt}}$. In Property A.6 of Appendix A, I show that these rents equal:

$$R_j^e = Y_j \times \left[\sum_{\chi \in \mathcal{X}} \int \left(\frac{1 + \alpha_j \boldsymbol{\varepsilon}_j(\chi, \varphi)}{1 + \boldsymbol{\varepsilon}_j(\chi, \varphi)} \right) \left(\frac{\theta_{j\chi} (L_j^{\text{eff}}(\chi))^{\rho_j}}{\sum_{\chi'} \theta_{j\chi'} (L_j^{\text{eff}}(\chi'))^{\rho_j}} \right) \left(\frac{\varphi L_j(\chi, \varphi)}{\int \varphi' L_j(\chi, \varphi') d\varphi'} \right) d\varphi - \alpha_j \left(\frac{Y_j^{\text{pt}}}{Y_j} \right) \right], \quad (17)$$

where Y_j^{pt} is the firm's value added as a price-taker. In this model, Y_j^{pt} does not appear to have a closed-form expression. Nevertheless, I show how to solve for this counterfactual quantity from the structural parameters in the model using a gradient descent algorithm. This procedure relies on the concavity of the firm's profit function, which is established in Appendix A.4.

In a model with homogeneous β , the scope for variation in firm rents is limited. As shown in Property A.6, this assumption implies that all firms must capture the same rents as a share of profit, R_j^e/Π_j , provided that they share the same returns to scale and that labor is perfectly substitutable.¹² Thus, for a simplified production technology (common $\alpha_j, \rho_j = 1$), all firms receive the same fraction of profits as rents, regardless of the differences in wages or amenities.

When β is heterogeneous, a firm's rents will depend on the composition of its workforce, which is shaped by the full set of firm characteristics, not just technology. This compositional channel makes the relationship between firm productivity T_j and rents R_j^e ambiguous. On the one hand, a more productive firm can capture higher profits. On the other hand, greater labor

¹²As the firm approaches constant returns to scale, $\alpha_j \rightarrow 0$, the profit that it could earn as a price-taker tends to zero. Therefore, any profit that the firm manages to capture would come in the form of rents: $\lim_{\alpha_j \rightarrow 0} \Pi_j = R_j^e$.

productivity raises wages at the firm, attracting workers with more wage-elastic labor supply. This reduces firm market power and makes labor more costly. Thus, while higher productivity can raise a firm's profit, it may also reduce the share of that profit that constitutes rents. In the limit, as $\boldsymbol{\varepsilon}_j(X) \rightarrow \infty$ for all X , a firm's profit converges to that of a price-taker, so $R_j^e \rightarrow 0$.

In my analysis, it is difficult to empirically distinguish between a firm's ex-ante and ex-post rents. This task would require detailed information about how firms choose and pay for amenities, which is not available in the administrative data. So, for identification, I assume that firms cannot adjust their amenities, at least during the time frame that the model is estimated. This assumption does not restrict how amenities correlate with a firm's productivity or technology. It also does not rule out the possibility that firms initially choose amenities to maximize profits.

Pass-through of TFP Shocks to Wages

Next, I examine the impact of a hypothetical shock to firm productivity on workers' wages, as measured by the elasticity $\frac{\partial \log W_j(X)}{\partial \log T_j}$. This object is commonly used to learn about rent sharing, as it reflects how an increase in a firm's value added would be passed on to workers in the form of higher pay.¹³ In Property A.7 of Appendix A, I show that $\left[\frac{\partial \log W_j(X)}{\partial \log T_j} \right]_X = (I - \mathbf{K}_j)^{-1} \mathbf{1}$, where:

$$\begin{aligned} \mathbf{K}_j[X, X'] = & \boldsymbol{\varepsilon}_j(\chi', \varphi') \left[(1 - \alpha_j - \rho_j) \left(\frac{\theta_{j\chi}(L_j^{\text{eff}}(\chi))^{\rho_j}}{\sum_{\chi''} \theta_{j\chi''}(L_j^{\text{eff}}(\chi''))^{\rho_j}} \right) \left(\frac{\varphi L_j(\chi, \varphi)}{\int \varphi'' L_j(\chi, \varphi'') d\varphi''} \right) \right. \\ & \left. - \mathbb{1}\{\chi = \chi'\} (1 - \rho_j) \frac{\varphi L_j(\chi, \varphi)}{\int \varphi'' L_j(\chi, \varphi'') d\varphi''} \right] + \mathbb{1}\{X = X'\} \frac{\frac{\partial^2 \log L_j(X)}{\partial [\log W_j(X)]^2}}{\boldsymbol{\varepsilon}_j(X)[1 + \boldsymbol{\varepsilon}_j(X)]}. \quad (18) \end{aligned}$$

To interpret the pass-through, first consider a case where β is homogenous. In this case, firms face a constant labor supply elasticity, so $\boldsymbol{\varepsilon}_j(X) = \boldsymbol{\varepsilon}$ and $\frac{\partial^2 \log L_j(X)}{\partial [\log W_j(X)]^2} = 0$ all j and X . As shown in Property A.7, this restriction implies that the pass-through can only vary due to differences in firm technology. In particular, if firms share the same returns to scale and labor is perfectly substitutable, then $\frac{\partial \log W_j(X)}{\partial \log T_j}$ is constant, which means that all workers, regardless of their firm or skills, experience the same proportional wage change in response to a firm productivity shock.

When β is heterogeneous, the pass-through changes in two ways. First, given that $\boldsymbol{\varepsilon}_j(X)$ is no longer constant, $\frac{\partial \log W_j(X)}{\partial \log T_j}$ depends on the full set of labor supply elasticities that firms face. This creates an additional source of variation within and across firms. Second, given how workers sort on individual heterogeneity, the wage markdown can amplify the impact of firm productivity on earnings. To see why, consider that when a productivity shock drives up wages at a firm, it shifts the composition of the firm's workforce toward workers with more elastic labor supply. This reduces the firm's wage-setting power and drives wages up further. In this way,

¹³For example, Card et al. (2018) explicitly interprets the pass-through $\frac{\partial \log W_j(X)}{\partial \log T_j}$ as a rent-sharing elasticity.

wage changes at a firm can be self-reinforcing.¹⁴ Note that, in principle, the pass-through could be greater than one, although my empirical results show that, in practice, it is well below one.

Welfare and Allocative Inefficiency

I consider a measure of social welfare that incorporates firm profits. Specifically, I assume that profits are redistributed to workers in direct proportion to their wages. This approach ensures that the redistribution does not distort workers' employment decisions. Welfare is defined as:

$$\mathcal{W} = E \left(\max_j \{u_{ij}(W_j(X_i), a_j(X_i))\} \right) + \log \left(\sum_{j=1}^J \Pi_j \right). \quad (19)$$

To evaluate the welfare cost of imperfect competition, I compare welfare in the monopsonistic economy to that which would be obtained by a social planner who optimally allocates labor to firms by solving: $\mathcal{W}^* = \max_{\{j(i)\}_i} \mathcal{W}$. I assume that the social planner has complete knowledge of worker and firm fundamentals but that firms continue to set wages based on their limited information of (β_i, ϵ_{ij}) . As I prove in Property A.8 of Appendix A, the planner's optimal allocation is equivalent to what would be achieved in a competitive (Walrasian) economy, in which firms act as price takers facing perfectly elastic labor supply curves. Moreover, I show that it is possible to solve for welfare under this optimal allocation from the model's structural parameters. In doing so, I can evaluate the welfare cost of imperfect competition in the market.

In an economy with constant markdowns, there would be no misallocation of labor across firms due to imperfect competition, because workers would sort into the same firms that they would choose in the absence of monopsony power. However, if markdowns vary across firms, then there is a potential for allocative inefficiency. This inefficiency occurs, for example, when a worker who is most productive at one firm j earns a higher wage at another firm k because firm j has a larger markdown than firm k . Such discrepancies can alter workers' rankings of firms and distort their labor supply decisions. Therefore, heterogeneous markdowns introduce a welfare cost of imperfect competition that is not present in the case of constant markdowns.¹⁵

¹⁴Another way to visualize this effect is to analyze the wage elasticity of the markdown $\frac{\epsilon_j(X)}{1+\epsilon_j(X)}$, which equals:

$$\frac{\partial \log \frac{\epsilon_j(X)}{1+\epsilon_j(X)}}{\partial \log W_j(X)} = \frac{\frac{\partial^2 \log L_j(X)}{\partial [\log W_j(X)]^2}}{\epsilon_j(X)[1+\epsilon_j(X)]}, \quad \text{where} \quad \frac{\partial^2 \log L_j(X)}{\partial [\log W_j(X)]^2} \geq 0.$$

This elasticity is positive, implying that wage changes at a firm are amplified through the markdown. In principle, this can generate multiple solutions to (12), potentially yielding two local optima—one with lower wages and employment but larger markdowns, and another with higher wages and employment but smaller markdowns. This phenomenon is discussed by Manning (2010) in the context of spatial agglomeration. In my setting, I show that multiplicity does not arise as long as $\text{Var}(\beta | X)$ is bounded relative to $1 - \alpha_j$. I also show that, even if multiple solutions to (12) exist, there is almost surely a unique global profit maximum (Appendix B.1, Lemma 2).

¹⁵Of course, with a nonemployment margin, the existence of any markdown can distort workers' labor-leisure

III. Data and Empirical Context

I now describe the data sources and the key variables that are used for my empirical analysis, as well as the criteria for sample selection. Additional details are provided in the Appendix.

III.A. Data Sources

For my analysis, I use a matched employer-employee panel dataset from Norway, containing information about workers and firms over the time period 1995-2018. The data is constructed by linking annual business tax records to a national employment register from the Norwegian Labour and Welfare Administration. It encompasses all workers who are not self-employed, and it covers the universe of firms that are required to file a tax return. Additional information about worker demographics and financial assets comes from individual tax filings and social security registers. All administrative data are maintained by Statistics Norway and are subject to comprehensive quality control measures. These data are considered to be highly reliable.¹⁶

At the worker-level, the data contains information about annual earnings, scheduled work hours, and the duration of a job spell. Earnings include a worker's salary, bonuses, and other payments that employers offer such as overtime, vacation, and severance pay.¹⁷ All monetary variables are adjusted for inflation and expressed in 2014 U.S. dollars. The data also reports a worker's education, experience, and specialization (e.g., college major and vocational training). These variables are used to define a worker's skill type χ . In the baseline specification, I define χ to be a combination of education and experience, where education is divided into three categories (no high school degree, high school degree but no college degree, and college degree) and experience is divided into two categories (less than 10 years in the workforce and more than 10 years in the workforce). This yields six distinct skill types. Alternate definitions of χ are considered in the Appendix. Lastly, the data contains rich demographic information about workers and their families. These data are used for making cross-sectional comparisons.

In the data, workers are linked to establishments, which represent sub-units (or branches) of a firm. To address this dimension, I measure employment at the establishment-level, and I assume that a firm's production function aggregates labor in a way that allows for imperfect substitutability between establishments.¹⁸ This approach allows for more flexible modeling

trade-offs, creating underemployment; yet, constant markdowns would still imply no misallocation across firms.

¹⁶See Fagereng et al. (2021) for discussion. Bruun et al. (2021) also give an independent quality evaluation.

¹⁷In-kind benefits are also included in earnings; however, parental leave and sickness benefits are excluded.

¹⁸Specifically, if the firm j operates multiple establishments $k \in \mathcal{K}_j$, then I extend the production function so that $Y_j = T_j [\sum_{k \in \mathcal{K}_j} \sum_{\chi \in \mathcal{X}} \theta_{jk\chi} (L_{jk}^{\text{eff}}(\chi))^{\rho_j}]^{(1-\alpha_j)/\rho_j}$, where $L_{jk}^{\text{eff}}(\chi) = \int \varphi D_{jk}(\chi, \varphi) d\varphi$ is the efficiency units of labor of skill type χ demanded by firm j for establishment k . This extension does not meaningfully affect the theoretical results or identification analysis. Thus, to ease notation, the ongoing discussion will continue to focus on a case where each firm operates only one establishment ($|\mathcal{K}_j| = 1$). As a robustness check in the Appendix, I assess the sensitivity of the parameter estimates to restricting the sample to single-establishment firms.

of sorting and wage determination. For example, if an oil and gas firm operates a processing plant in rural Norway and a corporate office in Oslo, then these two establishments are treated as distinct entities—both by workers and by the managing firm. Establishments are assigned to local labor markets based on industry and location. The Central Register of Establishments and Enterprises classifies them into 10 broad industries and 46 commuting zones, as defined in Bhuller (2009). I define a market as the intersection of an industry and a commuting zone.

Accounting data is measured at the firm-level and is collected from corporate income statements and balance sheets. A firm's value-added is defined as the difference between revenue and non-labor expenses, which include operating costs, spending on intermediate inputs, and capital depreciation. A firm's profit is defined as revenue minus overall costs, where overall costs account for wages paid to workers. All variables are measured prior to interest and taxes.

III.B. Sample Construction

Because my framework studies labor supply on the extensive margin, I restrict the estimation sample to full-time workers, defined as those who work at least 30 hours per week at a given establishment. Throughout my analysis, labor will be measured as full-time employment. To calculate earnings, I further restrict the sample to individuals aged 25 to 55 who work at the same establishment for an entire year and who do not have a higher-paying job elsewhere. Following Song et al. (2018), I also require that a worker's earnings are above some minimum threshold. In the baseline sample, I set this threshold to be \$15,000 per year. In the Appendix, I show that the empirical results are robust to changes in this minimum earnings threshold.

Part of my analysis relies on utilizing a balanced panel of firms. To achieve this, I restrict the estimation sample to firms that remain operational for nine consecutive years. During this period, the firm must consistently report positive revenue and maintain at least five full-time employees. In addition, the firm must continue operating the same set of establishments, and each establishment must stay in the same industry and commuting zone for the full nine years.

Table 1 provides descriptive statistics for the main estimation sample. More information about the distribution of wages, labor, value added, and profits can be found in the Appendix.

Table 1: Summary Statistics—Main Estimation Sample

	Unique Identifiers	Unique Observations
<i>Panel A. Population Counts</i>		
Number of Full-time Workers	1,325,863	8,303,075
Number of Establishments	31,252	475,614
Number of Firms	23,066	274,687
Number of Local Markets	457	10,563
	Share of Workforce	Mean Log Earnings
<i>Panel B. Worker Skill Types</i>		
Lower Secondary, <10 Years Exp.	6.87%	10.91
Upper Secondary, <10 Years Exp.	18.98%	11.11
College Graduate, <10 Years Exp.	14.74%	11.33
Lower Secondary, ≥ 10 Years Exp.	21.30%	11.09
Upper Secondary, ≥ 10 Years Exp.	23.67%	11.30
College Graduate, ≥ 10 Years Exp.	14.43%	11.59
	Mean	Std. Dev.
<i>Panel C. Accounting Variables</i>		
Log Revenues	18.05	2.20
Log Total Costs	18.01	2.19
Log Value Added	16.86	2.18
Log Profit	14.97	2.54

Notes. This table presents descriptive statistics for the baseline estimation sample. All monetary quantities are adjusted for inflation and are expressed in log 2014 US dollars.

IV. Identification of Economic Quantities

In this section, I demonstrate how to bring the model to data. In particular, I show how to recover key economic quantities of interest—such as the distribution of wage markdowns and rents, and measures of allocative inefficiency—from the matched employer-employee dataset. Some of these quantities can be recovered without full knowledge of the model’s structural parameters, but others—particularly those based on counterfactual outcomes—rely on the full set of parameters. For this reason, my goal is to achieve full point identification of the model.

My identification analysis follows two main steps. First, I show how to recover the firm-specific labor supply elasticities $\epsilon_j(X)$ by using panel data together with an economic shock that shifts the labor demand of multiple firms in the economy. Second, with these elasticities, I prove identification of firms’ production parameters $(T_j, \{\theta_{j\chi}\}_\chi, \rho_j, \alpha_j)$ and amenities $a_j(X)$, as well as the joint density $f_{\beta,X}$ that characterizes workers’ skills and wage-amenity trade-offs.

IV.A. Identification of Firm-Specific Labor Supply Elasticities

IV.A.1. Ideal Experiment

I begin by laying out the comparative statics that identify the labor supply elasticities. Consider a hypothetical, firm-specific productivity shock that shifts firm j 's TFP in logs from $\log T_j$ to $\log T_j + \delta_j$. Assume all other parameters in the economy remain unchanged. For every skill group X , let $(W_j(X), L_j(X))$ and $(W'_j(X), L'_j(X))$ represent the firm's equilibrium wage and employment before and after the shock, respectively. As $\delta_j \rightarrow 0$, these equilibrium outcomes can be used to recover the elasticity $\varepsilon_j(X)$. In particular, $\varepsilon_j(X)$ is defined by the condition:

$$\lim_{\delta_j \rightarrow 0} \left[(\log L'_j(X) - \log L_j(X)) - \varepsilon_j(X) \times (\log W'_j(X) - \log W_j(X)) \right] = 0. \quad (20)$$

This condition follows from two key properties of the model. First, under the equilibrium characterization, firms view each other to be strategically small. Therefore, they do not internalize wage changes that are confined to a single rival firm in the economy. This delivers an exclusion restriction, whereby an exogenous shock to T_j affects firm j 's labor supply $L_j(X)$ only through changes to its own wage $W_j(X)$. The second property of the model, established in Lemma 2 of Appendix B.1, is that the solution to the firm's first order condition is locally stable, which means that iterating on the equilibrium wage condition (11) will converge to a fixed point for nearby initial values. This property is important for the comparative statics because it ensures that a small, isolated productivity shock does not lead to large jumps in wages or employment.

Based on equation (20), an ideal experiment to recover the elasticities $\{\varepsilon_j(X)\}_{j,X}$ would be to introduce isolated, infinitesimal shocks to each firm's TFP and then measure the resulting change in each firm's wages and employment. This type of experiment is not feasible in practice, however, because I cannot independently manipulate each firm's productivity while holding all other aspects of the labor market fixed. Moreover, even if it were possible to do so, this would require measuring a firm's wages and employment conditional on each skill group $X = (\chi, \varphi)$. Since workers' skill levels φ are not observed in the data, the researcher would not be able to directly distinguish between a firm's labor inputs. In the remainder of this subsection, I show how to overcome these challenges using panel data and instrumental variables.

IV.A.2. Econometric Model

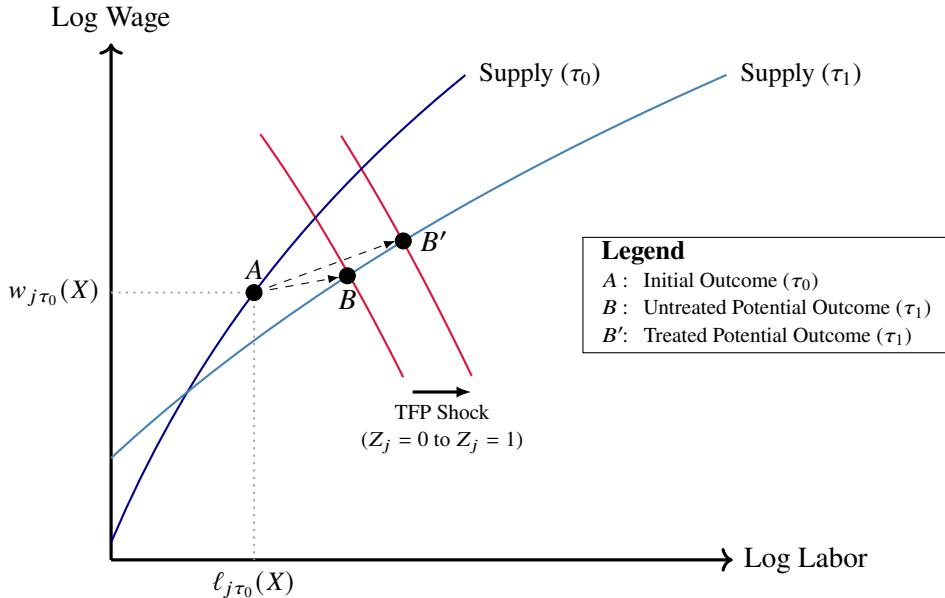
I now describe the empirical framework and assumptions that form the basis for my identification results. Throughout this discussion, I adopt the logit error structure, where I assume $F_\epsilon(\epsilon) = \exp(-\exp(-\epsilon))$. Under this specification, the firm-specific labor supply elasticities are defined by equation (10). To simplify notation, I use lowercase letters to denote variables that are expressed in logs, defining $w_j(X) = \log W_j(X)$, $\ell_j(X) = \log L_j(X)$, and $y_j = \log Y_j$.

Description of Empirical Framework for a Generic Instrument

Consider the evolution of equilibrium outcomes over two time periods: τ_0 and τ_1 . After the initial period τ_0 , an event occurs in which a subset of firms in the economy experiences a TFP shock, while the productivity of other firms remains unchanged. Following this event, wages and labor adjust to a new equilibrium by period τ_1 . Let $w_{j\tau_0}(X)$ and $\ell_{j\tau_0}(X)$ denote the log wages and employment for skill group X at firm j in period τ_0 , and let $w_{j\tau_1}(X)$ and $\ell_{j\tau_1}(X)$ be the corresponding outcomes in τ_1 . Define Z_j as a binary variable indicating whether firm j experiences a TFP shock. Firms are classified as *treated* if $Z_j = 1$ and *untreated* if $Z_j = 0$.

Figure 2 illustrates the evolution of wages and labor for a given firm j , delineating which aspects of the model remain constant and which vary over time. Importantly, between periods τ_0 and τ_1 , wages and labor adjust for all firms, including those that are untreated. This occurs because, unlike in the ideal experiment, the productivity shocks are not confined to just one firm. Rather, they take place at multiple firms in the economy. Consequently, wage adjustments among the treated firms would affect market aggregates by shifting workers' wage indices, $I(\beta, X)$. This, in turn, reshapes the labor supply curves faced by every firm, including untreated firms. Thus, the TFP shocks generate spillover effects that impact all firms in the economy.

Figure 2: Illustration of Potential Outcomes for a Firm



Notes. This figure depicts how a firm j 's potential wage and labor outcomes respond to an event that shifts TFP for multiple firms in the labor market. At τ_0 , workers' wage indices, skills and preferences are characterized by $\{I_{\tau_0}(\beta, X)\}_{\beta,X}$ and $\{F_{\beta,X|\tau_0}\}_X$, respectively, and at τ_1 , they are characterized by $\{I_{\tau_1}(\beta, X)\}_{\beta,X}$ and $\{F_{\beta,X|\tau_1}\}_X$, respectively. Apart from the TFP parameter T_j , which changes only for treated firms, all firm characteristics ($\{\theta_{j\chi}\}_\chi, \rho_j, \alpha_j, \{a_j(X)\}_X$) remain fixed between τ_0 and τ_1 .

Contributing to these dynamics, I allow workers' skills and preferences to evolve over time in response to the productivity shocks affecting firms. In particular, I permit the distribution of workers' skills F_X and the distributions $\{F_{\beta|X}\}_X$ governing workers' wage-amenity trade-offs to change from period τ_0 to τ_1 . I place no restrictions on the evolution of these distributions, imposing no constraints on how $(F_{X|\tau_0}, \{F_{\beta|X,\tau_0}\}_X)$ relates to $(F_{X|\tau_1}, \{F_{\beta|X,\tau_1}\}_X)$. While I also allow workers' idiosyncratic tastes to evolve over time, I assume that the cross-sectional distribution of these tastes retains a logit structure in each period, $F_\epsilon(\epsilon_{ij\tau}) = \exp(-\exp(-\epsilon_{ij\tau}))$. To achieve this, I assume that $\epsilon_{ij\tau}$ is a Markov process with independent innovations across i .

Given the evolution of workers' wage indices, skills, and preferences over time, firms face different labor supply curves in periods τ_0 and τ_1 . Therefore, in order to measure the elasticity of labor supply to any firm j , it is necessary to compare the firm's potential outcomes within the same time period. In period τ_0 , the treated and untreated potential outcomes are the same, since no TFP shocks have yet occurred. However, in period τ_1 , the potential outcomes differ, as illustrated by points B and B' in Figure 2. This difference reflects the comparative statics from the ideal experiment, capturing how a firm j 's wages and labor would respond to a TFP shock that is confined just to that firm. If it were possible to compare both potential outcomes for a firm in period τ_1 , then one could compute the firm's labor supply elasticity in that period.

Identification Assumptions

Throughout my identification analysis, I impose two key assumptions. First, I restrict how the exposure to a TFP shock, denoted by Z_j , depends on the inherent characteristics of a firm.

Assumption I (Instrument Independence). $Z_j \perp (\rho_j, \alpha_j, \{\theta_{j\chi}\}_\chi, \{a_j(X) - a_j(X')\}_{X,X'})$.

Assumption I states that the firm's treatment status Z_j is statistically independent of its technology (ρ_j, α_j) , the relative efficiencies $\{\theta_{j\chi}\}_\chi$ of each skill type, and the differences in amenities $\{a_j(X) - a_j(X')\}_{X,X'}$ between skill groups at the firm. Importantly, this assumption does not require Z_j to be fully exogenous. In particular, I allow it to depend on a firm's productivity through the baseline TFP T_j . Additionally, while Z_j is independent of amenity differences within a firm, it may still be affected by a firm's average amenity $E(a_j(X)|j)$. Consequently, treated and untreated firms can exhibit systematically different wages and labor.

Next, I make an assumption about the relationship between a worker's skill level φ and the other structural parameters in the model. This assumption consists of three main restrictions.

Assumption II (Restrictions on Unobserved Skill Levels).

- II.1.** $\varphi \perp \beta|\chi, \tau$.
- II.2.** $a_j(\chi, \varphi) = a_{j\chi} + a_{\chi\varphi}$.
- II.3.** $E(\log \varphi|\chi, \tau) = E(\log \varphi|\chi)$.

Assumption II.1 asserts that, given a worker's skill type χ and time period τ , preferences

over wages and amenities, as captured by the parameter β , do not depend on the skill level φ . This assumption does not imply that β is fully independent of worker skills, as I still allow β to depend on the skill type χ . Additionally, it does not rule out the possibility that β varies along unobserved skill dimensions; however, it does preclude systematic variation based on φ .

Assumption II.2 states that the non-wage amenity $a_j(\chi, \varphi)$ can be expressed as the sum of two components: one that is specific to the firm and the skill type, $a_{j\chi}$, and another that is specific to the skill level, $a_{\chi\varphi}$. This assumption ensures that amenities do not vary across firms with respect to unobserved aspects of worker skills. To better interpret the two amenity components $a_{j\chi}$ and $a_{\chi\varphi}$, I impose the normalization $E(a_{\chi\varphi}|\chi) = 0$ without loss of generality. This normalization allows me to define $a_{j\chi}$ as the average amenity at firm j for skill type χ .

Assumption II.3 requires that the average log skill level $E(\log \varphi|\chi, \tau)$ among workers in a given skill type χ is invariant over time τ . Note that this assumption does not restrict how the variance of these skill levels, or other higher-order moments, evolve. Moreover, it does not restrict the evolution of the distribution $F_{\chi|\tau}$, which governs workers' observable skill types. In my subsequent analysis, it will be useful to normalize φ by setting $E(\log \varphi|\chi) = 0$ for each $\chi \in \mathcal{X}$. This normalization is without loss of generality, as it is always possible to redefine the relative efficiencies $\{\theta_{j\chi}\}_{j,\chi}$ of each skill type to include the values of $\{E(\log \varphi|\chi)\}_\chi$.¹⁹

IV.A.3. Identification Results

I now provide results about the identification of firm-specific labor supply elasticities. First, I establish the identification of unobserved skill levels. Next, I outline the IV design used to recover the elasticities. Finally, I discuss the sources of instruments employed in my analysis.

Identification of Worker Skills

In order to draw inferences about the labor supply elasticities, I need to distinguish between workers with different skills. These skills are only partially observed in the data, however, as I do not directly measure a worker's skill level φ . To overcome this challenge, I rely on the restrictions implied by Assumption II. I demonstrate that, under this assumption, the skill levels φ are identified up to scale from observed differences in wages among workers at a given firm.

To establish identification, I first examine the implications of Assumption II for workers' labor supply decisions and the determination of wages at a firm. The key implication of this assumption is that the labor supply elasticities faced by a firm j do not vary along unobserved skill dimensions, i.e., $\boldsymbol{\varepsilon}_{j\tau}(\chi, \varphi) = \boldsymbol{\varepsilon}_{j\tau}(\chi)$. Thus, the wage markdowns at a firm j differ by workers' skill types χ , but not by their skill levels φ . This property implies the following result.

¹⁹Specifically, if $E(\log \varphi|\chi) \neq 0$ for some skill type χ , then I can redefine $\log \theta_{j\chi}$ as $\log \theta_{j\chi} + \rho_j E(\log \varphi|\chi)$.

Proposition 1. Under Assumption II, the wage can be written as: $W_{j\tau}(\chi_i, \varphi_i) = \varphi_{i\tau} \times \psi_{j\chi\tau}$, where the term $\psi_{j\chi\tau}$ is uniform across workers in a given firm j , skill type χ , and period τ .

This result provides a way to compute workers' skill levels from data about workers' earnings. In particular, I can recover φ from cross-sectional wage differences in each firm and skill type:

$$\log \varphi_{i\tau} = \log W_{j\tau}(\chi_i, \varphi_i) - E[\log W_{j\tau}(\chi_i, \varphi_i) | j, \chi, \tau]. \quad (21)$$

Going forward, it will be useful to measure wages in terms of efficiency units. I define the effective wage for workers with skill type χ at firm j in period τ as $W_{j\tau}^{\text{eff}}(\chi) = W_{j\tau}(\chi, \varphi)/\varphi$.²⁰ Using this notation, I can re-interpret the elasticity of labor supply to a firm as a function of the firm's effective wage. Specifically, I demonstrate that $\varepsilon_{j\tau}(\chi)$ takes the following form:

$$\varepsilon_{j\tau}(\chi) = \int \beta \times \frac{[\exp(\beta \log W_{j\tau}^{\text{eff}}(\chi)) / I_{\tau}^{\text{eff}}(\beta, \chi)] f_{\beta|\chi, \tau}(\beta|\chi, \tau)}{\int [\exp(\beta' \log W_{j\tau}^{\text{eff}}(\chi)) / I_{\tau}^{\text{eff}}(\beta', \chi)] f_{\beta|\chi, \tau}(\beta'|\chi, \tau) d\beta'} d\beta. \quad (22)$$

In this expression, $I_{\tau}^{\text{eff}}(\beta, \chi) = \sum_{k=1}^J \exp(\beta \log W_{k\tau}^{\text{eff}}(\chi) + a_{k\chi})$ can be interpreted as the effective wage index for workers with preference parameter β and skill type χ during time period τ .

Difference-in-Differences Estimand

I now present the main identification result of this subsection, which demonstrates how to recover the elasticity of labor supply to a firm using a difference-in-differences (DiD) research design. To motivate my approach, it is first important to recognize that I would not be able to determine the elasticities through either a cross-sectional comparison of treated and untreated firms or a comparison of outcomes within a firm over time. A cross-sectional comparison is not suitable because treated and untreated firms face different labor supply curves due to systematic variation in amenities. Moreover, a within-firm panel regression fails because of the spillover effects of TFP shocks on workers' wage indices and preferences. These spillovers imply that the same firm faces different labor supply curves in the pre- and post-periods.

I propose to recover the elasticity $\varepsilon_{j\tau_1}(\chi)$ from a ratio of DiD estimands, which control for a worker's skill type and a firm's effective wage in the pre-period. I define this estimand below.

Definition. Let $(\ell_{j\tau_0}(\chi), w_{j\tau_0}^{\text{eff}}(\chi))$ and $(\ell_{j\tau_1}(\chi), w_{j\tau_1}^{\text{eff}}(\chi))$ be the log labor and effective wage for skill type χ at firm j in periods τ_0 and τ_1 , respectively. The estimand $\text{DiD}_{\tau_0, \tau_1}(w|\chi)$ equals:

$$\frac{E[\ell_{j\tau_1}(\chi) - \ell_{j\tau_0}(\chi)|Z_j = 1, w_{j\tau_0}^{\text{eff}}(\chi) = w] - E[\ell_{j\tau_1}(\chi) - \ell_{j\tau_0}(\chi)|Z_j = 0, w_{j\tau_0}^{\text{eff}}(\chi) = w]}{E[w_{j\tau_1}^{\text{eff}}(\chi) - w_{j\tau_0}^{\text{eff}}(\chi)|Z_j = 1, w_{j\tau_0}^{\text{eff}}(\chi) = w] - E[w_{j\tau_1}^{\text{eff}}(\chi) - w_{j\tau_0}^{\text{eff}}(\chi)|Z_j = 0, w_{j\tau_0}^{\text{eff}}(\chi) = w]}.$$

²⁰A firm's effective wages can be recovered in logs as follows: $\log W_{j\tau}^{\text{eff}}(\chi) = E[\log W_{j\tau}(\chi, \varphi) | j, \chi, \tau]$.

The numerator of this expression is a DiD estimand corresponding to changes in log labor at a firm, while the denominator is a DiD estimand for changes in effective log earnings at a firm.

To show that this estimand corresponds to the labor supply elasticity, I first prove that the common trends assumption holds. Specifically, I show that—after controlling for a firm j 's initial effective wage $w_{j\tau_0}^{\text{eff}}(\chi)$ —the change in untreated potential outcomes over time, denoted by $w_{j\tau_1,0}^{\text{eff}}(\chi) - w_{j\tau_0,0}^{\text{eff}}(\chi)$ and $\ell_{j\tau_1,0}(\chi) - \ell_{j\tau_0,0}(\chi)$, can be written as a deterministic function of the firm-level structural parameters Γ_j , where $\Gamma_j = (\rho_j, \alpha_j, \{\theta_{j\chi}\}_\chi, \{a_j(X) - a_j(X')\}_{X,X'})'$. Assumption I ensures that a firm j 's treatment status Z_j does not depend on the parameters Γ_j . Therefore, under this assumption, treated and untreated firms with the same initial effective wages would, on average, experience the same change in their untreated potential outcomes.

Proposition 2. Suppose that Assumptions I and II hold. Then, given a firm's effective wage at τ_0 , the change in untreated potential outcomes is mean independent of the treatment status:

$$\begin{aligned} E\left[\ell_{j\tau_1,0}(\chi) - \ell_{j\tau_0,0}(\chi) \mid Z_j = 1, w_{j\tau_0}^{\text{eff}}(\chi) = w\right] &= E\left[\ell_{j\tau_1,0}(\chi) - \ell_{j\tau_0,0}(\chi) \mid Z_j = 0, w_{j\tau_0}^{\text{eff}}(\chi) = w\right]. \\ E\left[w_{j\tau_1,0}^{\text{eff}}(\chi) - w_{j\tau_0,0}^{\text{eff}}(\chi) \mid Z_j = 1, w_{j\tau_0}^{\text{eff}}(\chi) = w\right] &= E\left[w_{j\tau_1,0}^{\text{eff}}(\chi) - w_{j\tau_0,0}^{\text{eff}}(\chi) \mid Z_j = 0, w_{j\tau_0}^{\text{eff}}(\chi) = w\right]. \end{aligned}$$

Proposition 2 is key to my analysis. It implies that, by using a DiD research design, I can compare the treated and untreated potential outcomes for a firm in the post-period, thereby capturing the comparative statics from the ideal experiment. By this result, I demonstrate that the elasticity $\epsilon_{j,\tau_1}(\chi)$, as defined in equation (22), is point identified from the estimand $\text{DiD}_{\tau_0,\tau_1}(w_{j\tau_0}^{\text{eff}}(\chi)|\chi)$ for any firm j and skill type χ . Moreover, by evaluating these estimands over the effective wage distribution, I can recover the labor supply elasticities faced by each firm in the economy. As I discuss in Section V, this task relies on an *overlap condition*, asserting that $P(Z_j = 1 | w_{j\tau_0}^{\text{eff}}(\chi) = w) \in (0, 1)$ for all w in the support of $w_{j\tau_0}^{\text{eff}}(\chi)$. This condition guarantees that there are treated and untreated firms across the effective wage distribution.

Sources of Instruments

Up to this point, I have not specified how to construct the instrument Z_j . In my analysis, I consider both internal instruments—derived from existing data based on the model's implied restrictions—and external instruments, which draw on outside data sources. Both types of instruments may be used within my framework. However, they rely on different assumptions.

My main specification utilizes internal instruments, which are constructed by specifying a process by which the firm experiences productivity shocks. This approach follows Lamadon et al. (2022) by assuming that the firm's TFP evolves according to an autoregressive process, which has a unit root. Specifically, for each time period τ , I assume that log TFP is given by:

$$t_{j\tau} = \bar{t}_j + \tilde{t}_{j\tau}, \quad \text{where} \quad \tilde{t}_{j\tau} = \tilde{t}_{j\tau-1} + u_{j\tau}. \quad (23)$$

To ensure that the productivity shocks do not lead to large jumps in wages and labor, I assume that $u_{j\tau}$ is small relative to the log wage levels at firm j . I also make the following assumption.

Assumption III (Identification Assumptions for Internal Instruments).

III.1. (Instrument Relevance) $\text{Var}(u_{j\tau}) > 0$.

III.2. (Instrument Independence) $u_{j\tau} \perp (\rho_j, \alpha_j, \{\theta_{j\chi}\}_\chi, \{a_j(X) - a_j(X')\}_{X,X'})$.

Assumption III.1 asserts that productivity shocks vary across firms, which will imply that the internal instrument is relevant. Assumption III.1 guarantees that the internal instrument will satisfy Assumption I, which is sufficient to ensure that the common trends assumption holds.

Given this setup, I can construct the internal instruments from differences in the log value added of firms, $\Delta y_{j\tau}$, within a given calendar year τ . For integers $p, p' \geq 2$, I define the pre-period as $\tau_0 = \tau - p$, representing p years before τ , and the post-period as $\tau_1 = \tau + p'$, representing p' years after τ . Under Assumption III, I derive the following moment condition:

$$E \left[\Delta y_{j\tau} [(\ell_{j\tau_1}(\chi) - \ell_{j\tau_0}(\chi)) - \text{DiD}_{\tau_0, \tau_1}(w|\chi)(w_{j\tau_1}^{\text{eff}}(\chi) - w_{j\tau_0}^{\text{eff}}(\chi))] \mid w_{j\tau_0}^{\text{eff}}(\chi) = w \right] = 0. \quad (24)$$

This moment condition reflects an IV regression of long differences in log employment on long differences in log effective wages, instrumented by short differences in log value added, while controlling for effective wages in the pre-period. When implementing this procedure, I will define treatment status Z_j by separating firms at the median of the distribution of $\Delta y_{j\tau}$.²¹

To support the analyses using internal instruments, I also consider an external instrument. This instrument is based on plausibly exogenous product demand shocks arising from public procurement auction outcomes in Norway. This supplementary analysis follows the research design used by H. de Frahan et al. (2024), where a firm is classified as treated ($Z_j = 1$) in the year τ when it wins its first public procurement contract. Since this instrument is defined from outside data sources, it does not require making assumptions about the evolution of firm TFP. However, the main limitation of this instrument is that data on auction outcomes is only available for a subset of firms in the sample and for a limited time period, from 2003 to 2018.

IV.B. Identification of Technology, Amenities, and Preferences

I now demonstrate how to recover the structural parameters that characterize firm productivity, technology, and non-wage firm amenities, as well as the distribution of workers' preferences. These parameters are needed to determine rents for workers and firms, and they also allow me to run counterfactual analyses exploring sources of allocative inefficiency and inequality.

Firm Productivity and Technology

²¹In the Appendix, I show how to extend this framework to allow for transitory measurement error in value added and labor. I require that the errors are mutually independent and also independent of firm-level parameters.

I begin by proving identification of the production parameters $(T_{j\tau}, \{\theta_{j\chi}\}_\chi, \rho_j, \alpha_j)$ for any arbitrary firm j and any time period τ . Consistent with my previous analysis, I assume that the firm's TFP parameter $T_{j\tau}$ can evolve over time, while the parameters $(\{\theta_{j\chi}\}_\chi, \rho_j, \alpha_j)$ remain constant. In addition, I normalize the firm-specific efficiencies $\{\theta_{j\chi}\}_\chi$ by setting $\theta_{j\chi^*} = 1$ for some skill type $\chi^* \in \mathcal{X}$, as I can only recover these terms up to scale relative to the firm's TFP.

I demonstrate that identification relies on the ability to recover the labor supply elasticities faced by the firm at two different time periods. This task is feasible in my context because I have access to a 28-year panel, enabling me to replicate the analysis from Section IV.A over multiple years. For the time periods τ and τ' , respectively, I define $\varepsilon_{j\tau}(\chi)$ and $\varepsilon_{j\tau'}(\chi)$ as the labor supply elasticities for firm j corresponding to skill type χ . Using these elasticities, I can evaluate the marginal products of labor at the firm. Specifically, for $\tilde{\tau} \in \{\tau, \tau'\}$, I define:

$$\text{MPL}_{j\tilde{\tau}}^{\text{eff}}(\chi) = \frac{1 + \varepsilon_{j\tilde{\tau}}(\chi)}{\varepsilon_{j\tilde{\tau}}(\chi)} \times W_{j\tilde{\tau}}^{\text{eff}}(\chi) \quad (25)$$

as the effective marginal product for skill type χ at firm j , adjusting for workers' skill levels φ .

Using the structural wage equation (11), I can recover the substitution parameter ρ_j by analyzing how the relative effective marginal products and effective labor shares of different skill types change over time at a firm j . This result is formalized by the following proposition.

Proposition 3. For each firm j , the elasticity of substitution between different skill types is:

$$(1 - \rho_j)^{-1} = \frac{\log \left(\frac{L_{j\tau}^{\text{eff}}(\chi)}{L_{j\tau}^{\text{eff}}(\chi')} \right) - \log \left(\frac{L_{j\tau'}^{\text{eff}}(\chi)}{L_{j\tau'}^{\text{eff}}(\chi')} \right)}{\log \left(\frac{\text{MPL}_{j\tau}^{\text{eff}}(\chi)}{\text{MPL}_{j\tau}^{\text{eff}}(\chi')} \right) - \log \left(\frac{\text{MPL}_{j\tau'}^{\text{eff}}(\chi)}{\text{MPL}_{j\tau'}^{\text{eff}}(\chi')} \right)}, \quad \text{for any } \chi, \chi' \in \mathcal{X}.$$

Proposition 3 provides a new expression for the elasticity of substitution in terms of identified quantities in the model. This expression reflects how the relative demands for two labor inputs depend on their relative productivities, as measured by the elasticity of $L_{j\tau}^{\text{eff}}(\chi)/L_{j\tau}^{\text{eff}}(\chi')$ with respect to $\text{MPL}_{j\tau}^{\text{eff}}(\chi)/\text{MPL}_{j\tau}^{\text{eff}}(\chi')$. In my framework, this relationship is specific to each firm.

Once the substitution parameter ρ_j is recovered, I prove that $\{\theta_{j\chi}\}_\chi$, α_j , and T_j can be computed sequentially from the effective marginal product, labor, and value added at the firm.

Proposition 4. For any firm j , the parameters $\{\theta_{j\chi}\}_{\chi}$, α_j , and T_j are defined by the equations:

$$\begin{aligned}\theta_{j\chi} &= \exp \left[\log \left(\frac{\text{MPL}_{j\tau}^{\text{eff}}(\chi)}{\text{MPL}_{j\tau}^{\text{eff}}(\chi^*)} \right) + (1 - \rho_j) \log \left(\frac{L_{j\tau}^{\text{eff}}(\chi)}{L_{j\tau}^{\text{eff}}(\chi^*)} \right) \right]. \\ 1 - \alpha_j &= \exp \left[\log \text{MPL}_{j\tau}^{\text{eff}}(\chi) - y_{j\tau} - \log(\theta_{j\chi}) + (1 - \rho_j) \log L_{j\tau}^{\text{eff}}(\chi) + \log \sum_{\chi' \in \mathcal{X}} \theta_{j\chi'} \left(L_{j\tau}^{\text{eff}}(\chi') \right)^{\rho_j} \right]. \\ T_{j\tau} &= \exp \left[y_{j\tau} - \frac{1 - \alpha_j}{\rho_j} \log \sum_{\chi \in \mathcal{X}} \theta_{j\chi} \left(L_{j\tau}^{\text{eff}}(\chi) \right)^{\rho_j} \right].\end{aligned}$$

Note that recovering the substitution parameter ρ_j is the only step that requires knowledge of labor supply elasticities for two different time periods. If this parameter is already identified, or if one assumes that there is perfect substitutability between skill types (i.e., $\rho_j = 1$), then the other parameters $\{\theta_{j\chi}\}_{\chi}$, α_j , and T_j can be recovered from a single time period of data.

Firm Amenities

Next, I prove identification of the parameters $\{a_{j\chi}\}_{j,\chi}$, which represent the average amenities at each firm j for any skill type χ . In my analysis, these terms are only identified up to scale, as they can be shifted in a way that does not impact workers' labor supply decisions. Therefore, I will normalize the amenities by defining a reference firm j^* such that $a_{j^*\chi} = 0$ for all $\chi \in \mathcal{X}$.

To recover firms' amenities, I use a revealed preference approach. This method is based on the premise that, after accounting for the differences in workers' wages, a firm that has superior amenities for a given skill type χ would attract more workers of that type. Consequently, by comparing the employment levels among firms in a counterfactual setting where all firms offer the same wages, I could draw inference about the differences in firms' amenities.

To make these comparisons, I utilize the labor supply elasticity curves $\boldsymbol{\varepsilon}_{j\tau}(\chi, w)$, which are identified for each skill type χ by the estimand $\text{DiD}_{\tau_0, \tau_1}(w|\chi)$. By integrating along these curves, I can infer the labor supply to any firm j if it were to offer a different wage, holding fixed all other wages in the economy. In doing so, I can evaluate the amenities at firm j by comparing j 's labor supply to that of the reference firm j^* in the case where j and j^* pay the same wages. To illustrate this identification strategy, consider the following proposition.

Proposition 5. Under Assumptions I and II, the average amenity for skill type χ at firm j is:

$$a_{j\chi} = \ell_{j\tau}(\chi) + \int_{w_{j\tau}^{\text{eff}}(\chi)}^{w_{j^*\tau}^{\text{eff}}(\chi)} \boldsymbol{\varepsilon}_{j\tau}(\chi, w) dw - \ell_{j^*\tau}(\chi).$$

Worker Preferences

I conclude this section by establishing the identification of the densities $\{f_{\beta|\chi,\tau}\}_{\chi}$ that characterize the distribution of workers' marginal rates of substitution between wages and amenities. This identification is achieved through revealed preference, establishing a direct link between the distribution of workers' preferences and the shape of the labor supply curve faced by a firm.

To understand the source of identification, consider how the distribution of β influences a firm's labor supply curve. When β is homogeneous, this curve is isoelastic, and the elasticity of labor supply is equal to β . However, when β is heterogeneous, the curve is nonlinear, and the labor supply elasticity depends on the firm's wage. As shown by equations (10) and (11), the nature of this dependence carries information about the moments of the distribution of β .

In the following proposition, I demonstrate that this relationship extends even further. In particular, the elasticity curve $\varepsilon_{j\tau}(\chi, w)$ uniquely corresponds to the density $f_{\beta|\chi,\tau}(\beta|\chi, \tau)$.

Proposition 6. Given wages and labor $\{W_{j\tau}^{\text{eff}}(\chi), L_{j\tau}(\chi)\}_{j=1}^J$, there is a one-to-one mapping between the labor supply elasticity curve $\varepsilon_{j\tau}(\chi, w)$ and the density function $f_{\beta|\chi,\tau}(\beta|\chi, \tau)$.

Proposition 6 implies that the density $f_{\beta|\chi,\tau}$ is point identified under same the assumptions that are needed to recover the labor supply elasticity curve $\varepsilon_{j\tau}(\chi, w)$. The proof leverages the fact that a firm's labor supply curve, as characterized by (10), can be expressed as a Laplace transform. Moreover, I show that the inverse of this transform is proportional to the density of β . This property establishes the invertibility of a firm's labor supply curve to obtain $f_{\beta|\chi,\tau}$.²²

V. Estimation Strategy and Empirical Results

In this section, I describe my estimation procedure and present my empirical results, including estimates of firm-specific labor supply elasticities and structural parameters. My results show significant heterogeneity in the labor supply elasticities, indicating that wage markdowns for workers vary considerably both within and across firms. Additionally, I find that firms offering lower wages to workers in a given skill group are more likely to face smaller elasticities. This pattern aligns with the model's predictions, suggesting that firms with lower effective wages tend to exert greater wage-setting power. When comparing across skill groups, I find that the labor supply to firms is generally less elastic for more experienced workers. Finally, my estimates indicate that there is substantial variation in firm productivity, technology, and non-wage amenities, suggesting that these factors all contribute to systematic worker sorting.

V.A. Estimation of Firm-Specific Labor Supply Elasticities

This subsection addresses the estimation of labor supply elasticities. I begin by describing my methodology and my specification choices, followed by a discussion of the empirical findings.

²²Any two functions share the same Laplace transform only if they differ on a set of Lebesgue measure zero.

Estimation Procedure and Specifications

My estimation procedure is based on the identification strategy outlined in Section IV.A. This approach consists of two key steps. First, I estimate workers' skill levels by implementing a two-way fixed effects estimator, regressing log wages on a worker fixed effect γ_i and a fixed effect $\psi_{\chi(i,\tau),j(i,\tau),\tau}$ for the skill type, firm, and year. From this regression, I define a worker's skill level φ_i as $\exp(\gamma_i)$ and the effective wage $W_{j\chi\tau}^{\text{eff}}$ as $\exp(\psi_{j(i,\tau),\chi(i,\tau),\tau})$. Table 2 presents basic summary statistics on workers' skill levels, with more details provided in the Appendix.

Next, I estimate the labor supply elasticities for each firm by computing a sample analogue of the estimand $\text{DiD}_{\tau_0,\tau_1}(w|\chi)$. This estimand corresponds to a ratio of DiD estimands, controlling for a firm's effective wage in the pre-period. Importantly, because there is a finite number of firms in the data, it is unlikely that I would observe multiple firms offering exactly the same effective wages. To overcome this issue, I employ a nonparametric Kernel estimator:

$$\frac{\sum_j K_{1,j}(w) \mathbf{1}\{Z_j = 1\} (\ell_{j\tau_1}(\chi) - \ell_{j\tau_0}(\chi)) - \sum_j K_{0,j}(w) \mathbf{1}\{Z_j = 0\} (\ell_{j\tau_1}(\chi) - \ell_{j\tau_0}(\chi))}{\sum_j K_{1,j}(w) \mathbf{1}\{Z_j = 1\} (w_{j\tau_1}^{\text{eff}}(\chi) - w_{j\tau_0}^{\text{eff}}(\chi)) - \sum_j K_{0,j}(w) \mathbf{1}\{Z_j = 0\} (w_{j\tau_1}^{\text{eff}}(\chi) - w_{j\tau_0}^{\text{eff}}(\chi))}, \quad (26)$$

where $K_{1,j}(w)$ and $K_{0,j}(w)$ denote the Kernel weights corresponding to treated and untreated populations, respectively. This estimator allows me to compare the outcomes of firms with similar, though not identical, initial wages, thereby ensuring that the method is tractable for finite sample sizes. In Appendix B.8, I prove that this estimator is consistent. Additionally, I derive the asymptotic properties of the estimator, demonstrating that it is capable of inference.

In my baseline specification, I utilize a Gaussian kernel function to construct the weights $K_{z,j}(w)$.²³ As I show in the Appendix, the point estimates are robust to alternative kernel functions, including the Uniform kernel, which corresponds to a standard nonparametric binning estimator. I also investigate the sensitivity of the estimates to the bandwidth choice, finding that the results remain relatively stable across varying bandwidths. Collectively, these robustness analyses provide confidence in the estimator's ability to approximate $\text{DiD}_{\tau_0,\tau_1}(w|\chi)$.

My primary estimates are based on internal instruments, following the same identification arguments laid out in Section IV.A. For each year τ that TFP shocks occur, I define treatment status Z_j by separating firms at the median of the distribution of $\Delta y_{j\tau}$. I then define the pre- and post-periods to be $\tau_0 = \tau - 2$ and $\tau_1 = \tau + 3$, respectively. I estimate the model repeatedly using all years of available data, with the post-period τ_1 spanning from 2002 to 2017. As a robustness check, I also calibrate the model using estimates obtained with external instruments based on public procurement contract wins in Norway. These calibration results draw on estimates from H. de Frahan et al. (2024), who analyze the same Norwegian admin-

²³Specifically, I set $K_{z,j}(w) = \frac{1}{\sqrt{2\pi}} \exp\left[-\frac{1}{2}\left(\frac{1}{h}(w_{j\tau_0}^{\text{eff}}(\chi) - w)\right)^2\right]$ where the parameter h governs the bandwidth.

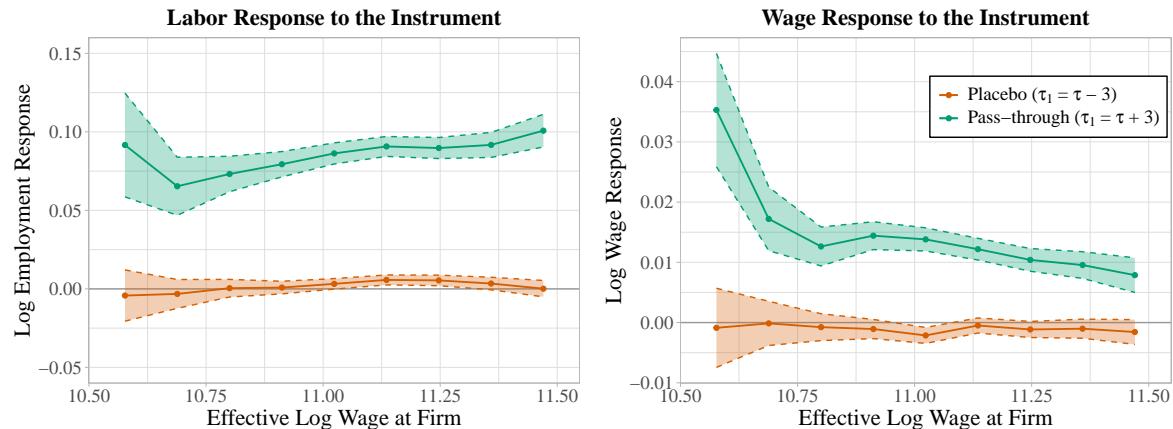
istrative dataset. I present my findings in the Appendix, along with a description of the data sources used in the original paper. Both internal and external instruments yield similar results.

When implementing my approach, I include fixed effects for local labor markets, which represent the interactions between industries and commuting zones as defined in Section III. For comparison, I also estimate the model without fixed effects, treating Norway as one single labor market. Those results can be found in the Appendix. I average my estimates across all available years of data. Since my framework allows the labor supply elasticities to vary over time, my estimates may be interpreted as an average of these quantities across multiple years.

Description of Empirical Findings

Figure 3 presents the reduced form estimates, averaged across workers' skill types and plotted over the distribution of effective wages. These estimates correspond to the numerator and denominator of the DiD estimator that is defined in equation (26). In this figure, the green line displays estimates for $\tau_1 = \tau + 3$, while the orange line depicts the difference in pre-period trends between treated and untreated firms, with τ_1 set to $\tau - 3$. Notably, the placebo effect does not significantly differ from zero, lending further support to the validity of the IV design.

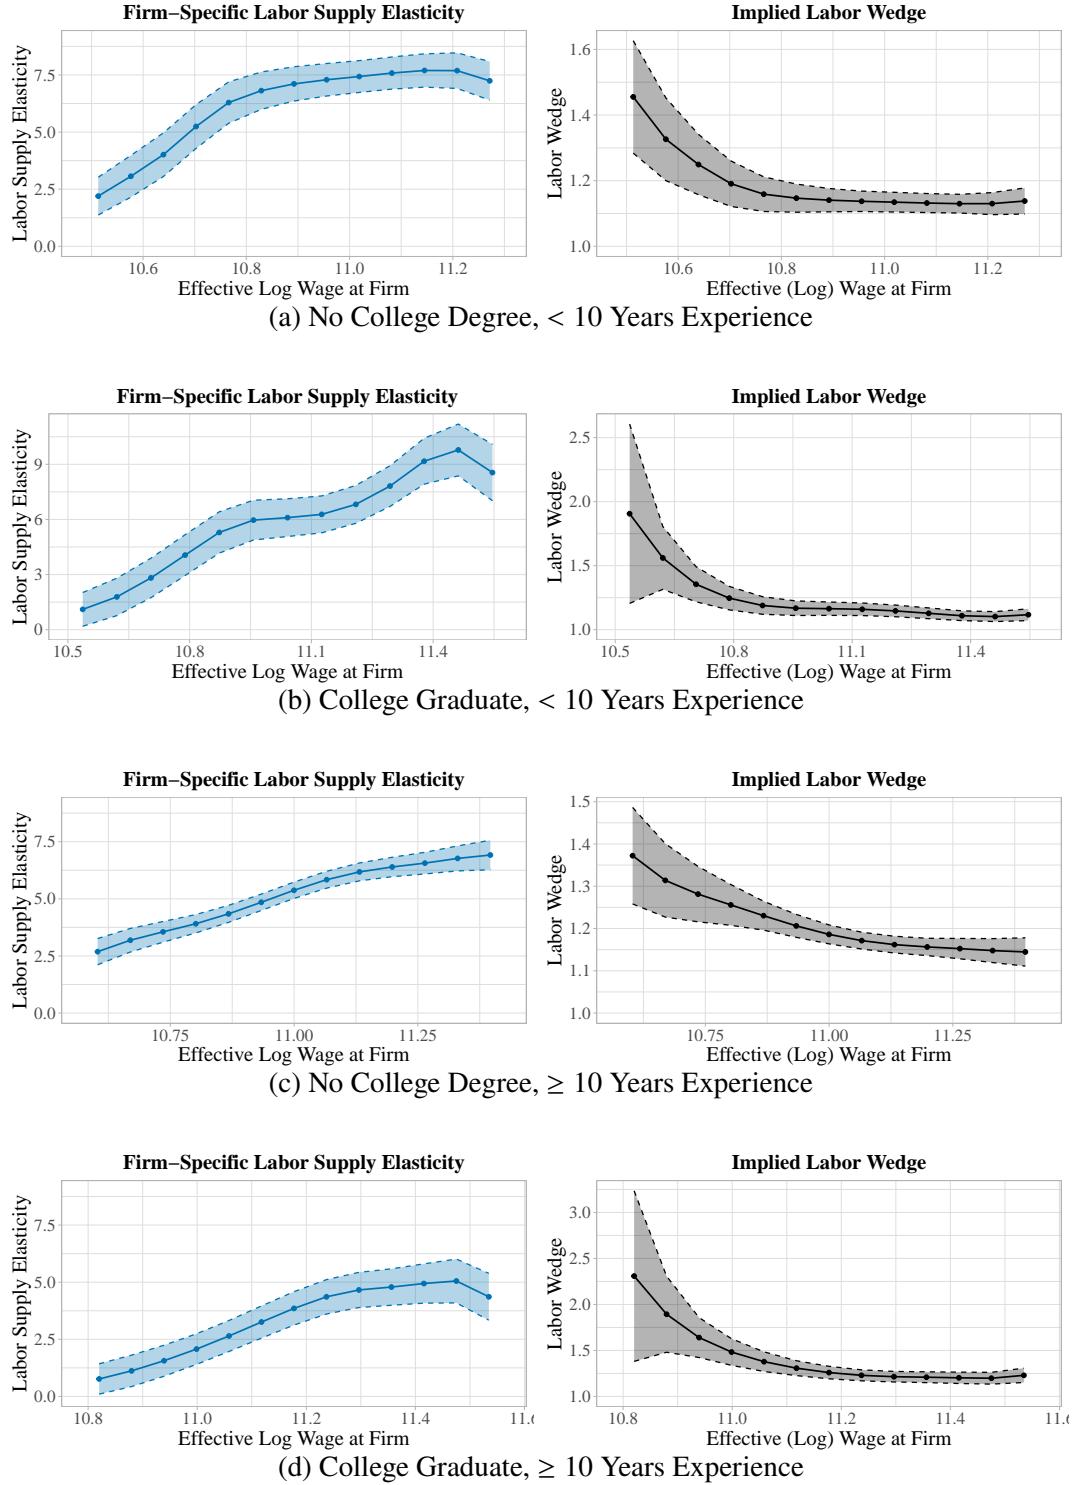
Figure 3: Average Pass-Through Rates by Effective Wage



Notes. This figure plots the reduced form estimates, averaged across skill types and years $\tau_1 \in \{2002, \dots, 2017\}$. This specification includes market fixed effects. Standard errors are bootstrapped using 500 bootstrap samples.

The reduced form estimates reveal two distinct patterns. First, the log labor response to the instrument consistently ranges between 5% and 10%, suggesting that employment growth is relatively independent of firm size. This observation aligns with Gibrat's law. Second, the log wage response ranges from 1% to 4%, and it decreases as a firm's effective wage rises. Based on equation (18), this pattern may suggest that higher-paying firms have lower returns to scale. Additionally, it could indicate that these firms face larger labor supply elasticities.

Figure 4: IV Estimates of Labor Supply Elasticities and Labor Wedges



Notes. This figure plots IV estimates of firm-specific labor supply elasticities and labor wedges. These estimates are averaged across all years of available data, with $\tau_1 \in \{2002, \dots, 2017\}$. The subgroup “No College Degree” represents a worker-weighted average of two different skill types: “Lower Secondary” and “Upper Secondary”. This specification includes market fixed effects. Standard errors are bootstrapped using 500 bootstrap samples.

Figure 4 presents the IV estimates of firm-specific labor supply elasticities, disaggregated by workers' college attainment and experience. The figure also plots the implied labor wedge at each firm, defined as the ratio of the marginal product to a worker's wage. These estimates reveal that the labor supply elasticities differ substantially across firms, ranging between 1.5 and 10.0 for certain skill types. Moreover, they indicate that firms with higher effective wages face larger elasticities. This pattern is generally similar across worker skill types, and it implies that market power is more concentrated at the lower end of the effective wage distribution.

Although the estimated labor supply elasticities are broadly consistent across skill types, I observe notable differences based on a worker's experience. Specifically, I find that workers who have less job experience tend to exhibit larger labor supply elasticities. In addition, these elasticities are considerably more variable across firms. Under my model, these differences imply that less experienced workers are more likely to prioritize higher wages when choosing a job, and that there is also more dispersion in wage-amenity trade-offs among these workers.

To further validate my approach, I conduct two additional robustness checks. First, I test whether the instrument Z_j impacts labor on the intensive margin by influencing work hours. This test addresses a potential concern that firms respond to productivity shocks by adjusting workers' scheduled hours instead of their wages. If this were the case, then it would threaten the validity of my IV design, which assumes that the entire pass-through of demand shocks occurs in the form of wage adjustments, not changes in work hours. I find no evidence that these shocks significantly affect workers' hours, which reinforces the credibility of my assumptions.

Second, I test the assumption that firms are strategically small in the economy. A testable implication of my model is that firms offering identical wages face the same labor supply elasticities. However, this property need not hold if firms are strategically large. In a concentrated labor market, firms that offer the same wages could still face different labor supply elasticities if they have different market shares. Given this property, I test whether a firm's labor supply elasticity varies with respect to the market share, after controlling for the effective wages. I find no evidence of a market share effect, which lends support to the validity of my framework.

V.B. Estimation of Structural Parameters in the Model

Next, I present estimates of the structural parameters in the model, including the production parameters $(T_{j\tau}, \{\theta_{j\chi}\}_\chi, \rho_j, \alpha_j)$, the amenities $a_{j\chi}$, and the densities $f_{\beta|\chi,\tau}$ that characterize workers' preferences. These estimates are summarized in Table 2, as well as Figures 5 and 6.

Firm Productivity, Technology, and Amenities

To estimate the production parameters, I apply the formulas provided in Propositions 3 and 4, which rely on the elasticity estimates, the effective wages, and the value added for each firm. I perform these calculations separately for every firm j and year τ , giving me a unique set of

estimates for each combination of j and τ in the sample.²⁴ In Figure 5, I plot the distributions of firms' estimated TFPs $T_{j\tau}$, returns to scale $1 - \alpha_j$, and substitution parameters ρ_j , averaged across all years of data. The means and variances of these quantities are reported in Table 2. I observe significant variation in firms' returns to scale and substitution parameters, suggesting that differences in technology across firms may play an important role in wage determination.

Table 2: Estimates of Structural Parameters—Means and Variances

	Notation	Mean	Variance
<i>Panel A. Firm Production</i>			
Log Total Factor Productivity	$t_{j\tau}$	12.92 (0.05)	1.19 (0.01)
Returns to Scale	$1 - \alpha_j$	0.70 (0.01)	0.04 (0.01)
Substitution Parameter	ρ_j	0.72 (0.04)	0.54 (0.13)
<i>Panel B. Firm Amenities</i>			
	$a_{j\chi}$		
No College, < 10 Years Exp.	—	1.54 (0.06)	
College, < 10 Years Exp.	—	2.02 (0.02)	
No College, \geq 10 Years Exp.	—	1.63 (0.11)	
College, \geq 10 Years Exp.	—	1.60 (0.04)	
<i>Panel C. Worker Skill Levels</i>			
	φ_i		
No College, < 10 Years Exp.	—	0.062 (0.01)	
College, < 10 Years Exp.	—	0.073 (0.02)	
No College, \geq 10 Years Exp.	—	0.064 (0.00)	
College, \geq 10 Years Exp.	—	0.076 (0.01)	
<i>Panel D. Worker Preferences</i>			
	β_i		
No College, < 10 Years Exp.	6.30 (0.76)	13.76 (1.01)	
College, < 10 Years Exp.	7.23 (1.05)	16.25 (2.71)	
No College, \geq 10 Years Exp.	5.96 (0.42)	15.71 (0.87)	
College, \geq 10 Years Exp.	4.28 (0.68)	15.11 (2.25)	

Notes. This table presents estimates of cross-sectional means and variances of model parameters, averaged across years. The means are omitted for parameters that are only identified up-to-scale.

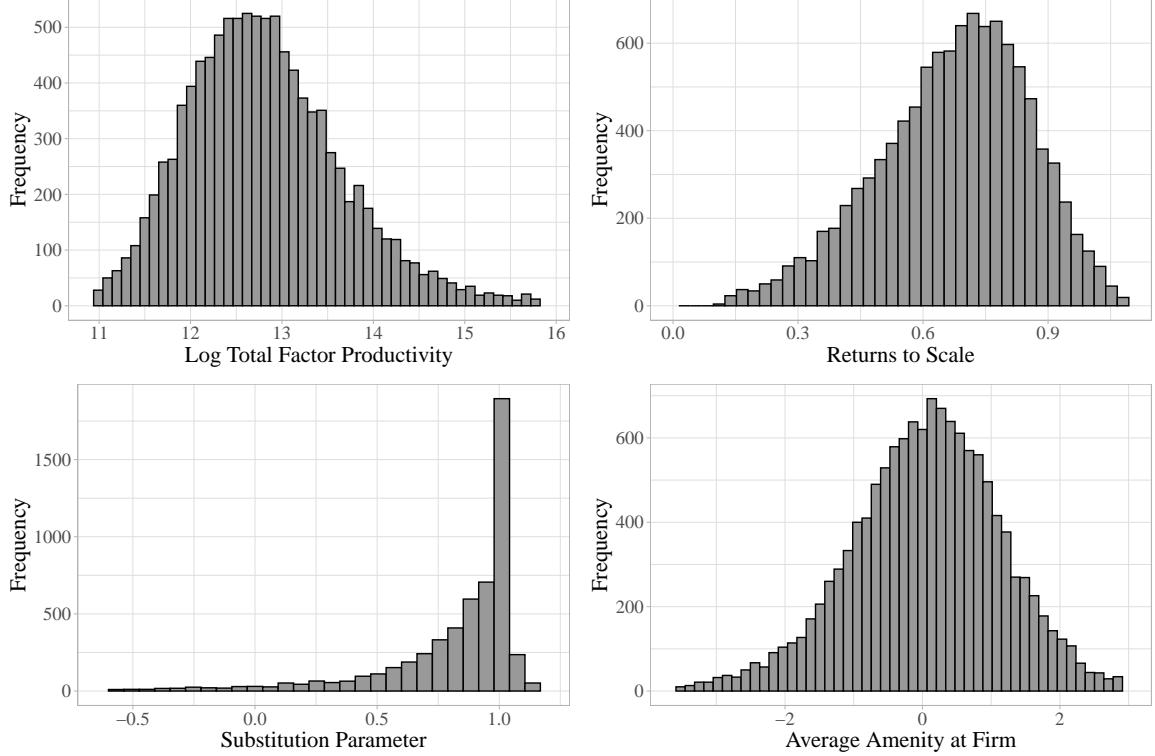
To estimate non-wage amenities, I follow the procedure laid out in Proposition 5, which involves integrating the estimated labor supply elasticities with respect to a firm's effective wage.²⁵ This approach allows me to compare labor across firms in a counterfactual setting where firms offer identical wages. I plot the estimated distribution of firms' average amenities in Figure 5, with variances by subtype shown in Table 2. I find that the amount of variation

²⁴Recall that computing the substitution parameter ρ_j for a firm requires me to compare estimates over two different time periods: τ and τ' . In my implementation, I define τ' as $\tau - 1$, corresponding to the previous year.

²⁵For each skill type χ , amenities are normalized relative to a “reference firm” j^* , where $a_{j^*\chi} = 0$. In my analysis, j^* is defined as the firm offering an average wage equal to the median of the distribution in the sample.

in amenities is generally similar across worker types. However, for less experienced workers, it appears that those with a college education face greater variation in amenities across firms.

Figure 5: Histograms of Estimated Firm Productivity, Technology, and Amenities



Notes. This figure plots histograms of the estimated (log) total factor productivity $\log T_{j\tau}$, the returns to scale $1 - \alpha_j$, the substitution parameter ρ_j , and the average amenities $E(a_{j\chi}|j)$ across firms in the estimation sample.

Worker Preferences

To estimate the density function $f_{\beta|\chi,\tau}$, I draw on Proposition 6. This result establishes that the density is uniquely determined from the labor supply curve faced by an individual firm. Consequently, $f_{\beta|\chi,\tau}$ may be recovered as the unique solution to the following equations:

$$L_{j\tau}(\chi, W) = \int \frac{\exp(\beta \log W + a_{j\chi})}{I_\tau^{\text{eff}}(\beta, \chi)} f_{\beta|\chi,\tau}(\beta|\chi, \tau) f_{\chi|\tau}(\chi|\tau) d\beta, \quad \text{for } W > 0. \quad (27)$$

Note that the labor supply curve $L_{j\tau}(\chi, W)$, the effective wage indices, $I_\tau^{\text{eff}}(\beta, \chi)$, the density of skill types $f_{\chi|\tau}(\chi|\tau)$, and the non-wage amenities $a_{j\chi}$ can all be recovered from the data.

To make the estimation problem tractable, I introduce two important modifications. First, given that the density function $f_{\beta|\chi,\tau}$ is an infinite-dimensional parameter, I will approximate

it using a finite-dimensional basis expansion. Specifically, I replace $f_{\beta|\chi,\tau}$ with $\hat{f}_{\beta|\chi,\tau}$, where:

$$\hat{f}_{\beta|\chi,\tau}(b) = \sum_{k=1}^{d_\phi} \phi_k \hat{f}_{\beta|\chi,\tau}^{(k)}(b), \quad \text{for } \phi = (\phi_1, \dots, \phi_{d_\phi}) \in \Phi. \quad (28)$$

In this expression, ϕ is a d_ϕ -dimensional vector of unknown coefficients with support Φ , and $\{\hat{f}_{\beta|\chi,\tau}^{(k)}\}_k$ are basis functions known to the researcher. By defining the basis in this way, I can approximate $f_{\beta|\chi,\tau}$ with a function $\hat{f}_{\beta|\chi,\tau}$ that is parametrized by a finite-dimensional vector ϕ .

Second, rather than imposing an infinite number of constraints, I evaluate equation (27) on a finite grid of effective wages $\mathbf{W} \subset \mathbb{R}_+$, noting that identification is achieved as $|\mathbf{W}| \rightarrow \infty$. This modification allows me to recast (27) as a convex optimization problem, minimizing the ℓ^2 -distance between the estimated choice probabilities and the predicted probabilities under the density $\hat{f}_{\beta|\chi,\tau}$. In particular, my estimation procedure involves choosing ϕ^* by solving:

$$\phi^* = \operatorname{argmin}_\phi \sum_{W \in \mathbf{W}} [\phi' \hat{h}_\tau(\chi, W) - \hat{L}_{j\tau}(\chi, W)]^2. \quad (29)$$

where $\hat{L}_{j\tau}(\chi, W)$ is the estimated labor supply for $W_{j\tau}^{\text{eff}}(\chi) = W$ and $\hat{h}_\tau(\chi, W) \in \mathbb{R}^{d_\phi}$ is a vector with entries $\hat{h}_{\tau,k}(\chi, W) = \int [\exp(\beta \log W + \hat{a}_{j\chi}) / \hat{I}_\tau^{\text{eff}}(\beta, \chi)] \hat{f}_{\beta|\chi,\tau}^{(k)}(\beta|\chi, \tau) \hat{f}_{\chi|\tau}(\chi|\tau) d\beta$.²⁶

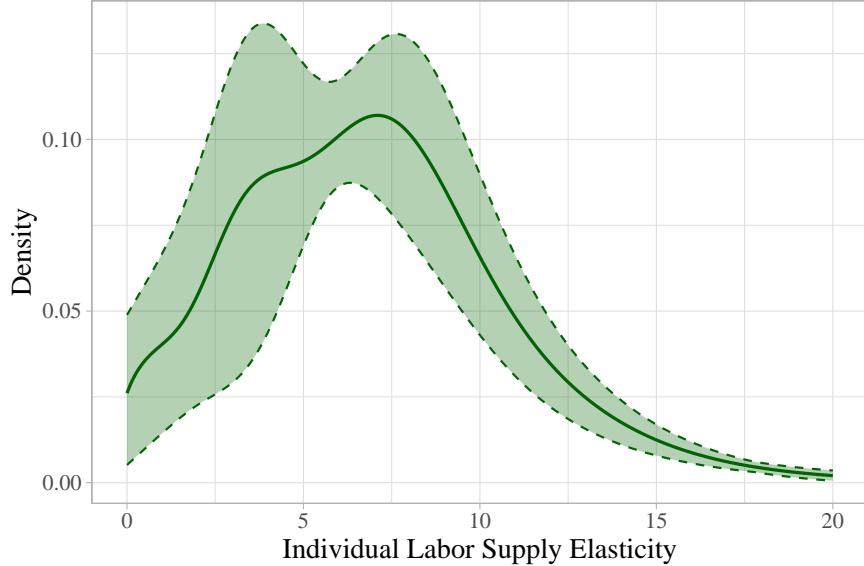
In my baseline specification, I construct a basis using mixtures of $d_\phi = 18$ Gamma density functions. This type of density have been shown to be particularly effective for nonparametric estimation of functions that have positive support; e.g., see Wiper et al. (2001) for discussion. As a robustness check, I investigate the sensitivity of my resulting estimates to the choice of basis functions, the basis size d_ϕ , and the number of points $|\mathbf{W}|$ in the wage grid.²⁷ I find that the estimates remain robust to variation in each of these elements. Moreover, I demonstrate that the earnings and employment distributions predicted by the estimated densities $\hat{f}_{\beta|\chi,\tau}$ closely match those in the sample, which suggests that my estimates are a good fit for the data.

Figure 6 plots the estimated density of β among workers in the population. I also compare moments of this distribution for different worker types in Table 2. I find that more experienced workers tend to have lower values of β , which is associated with a greater willingness to pay for amenities. However, there is substantial variation in preferences within each worker type.

²⁶Note that equation (29) may be written in matrix form as $\phi^* = \operatorname{argmin}_\phi \{\phi^\top \hat{h}^\top \hat{h}\phi - 2\phi^\top \hat{h}^\top \hat{L} + \hat{L}^\top \hat{L}\}$, where $\hat{L} \in \mathbb{R}^{|\mathbf{W}|}$ and $\hat{h} \in \mathbb{R}^{|\mathbf{W}| \times d_\phi}$. This framing of the problem is more convenient for computational implementation.

²⁷For example, I also approximate $f_{\beta|\chi,\tau}$ using a linear sieve constructed from Laguerre polynomials. These functions are capable of approximating any square-integrable function with positive support (see Chen, 2007).

Figure 6: Density of Workers’ Individual Labor Supply Elasticities



Notes. This figure plots the estimated density f_β , which corresponds to the distribution of worker’s preferences over wages and amenities. The confidence region is bootstrapped using estimates from 500 bootstrap samples.

VI. Implications for Rent Sharing, Misallocation, and Inequality

I now summarize the main empirical findings from my estimates, quantifying imperfect competition in the labor market and assessing its implications for welfare and earnings inequality.

VI.A. Quantifying Imperfect Competition

Table 3 presents the estimates of wage markdowns, along with worker and employer rents. On average, I find that workers are paid approximately 86% of their marginal product. However, these markdowns vary both within and across firms, ranging from approximately 70% to 95% at the 1st and 99th percentiles. I estimate that about 60% of the variation in markdowns takes place within firms, while the remaining 40% is due to differences between firms. The variation within firms is explained by systematic differences in workers’ wage-amenity trade-offs across different skill types χ . Meanwhile, the variation between firms arises from differences in wage-amenity trade-offs for equally-skilled workers. Note that this variation would be overlooked without accounting for unobserved heterogeneity in workers’ wage-amenity trade-offs.

I find that workers and firms capture substantial rents, with worker rents being larger, on average, than firm rents. I estimate that the average worker is willing to pay approximately \$13,500—or 17% of their earnings—to avoid working for a different firm. In contrast, the average firm obtains about \$9,400 per worker—or 29% of its profit—by exerting wage-setting power. Figure A.10 of the Appendix plots the cross-sectional densities of estimated worker and firm rents, averaged across years. I find that these rents vary considerably in the economy.

Table 3: Estimates of Wage Markdowns and Rents

	Mean	Std. Dev.	Decomposition of Variation	
			Within Firm	Between Firm
<i>Wage Markdown</i>	0.864 (0.011)	0.035 (0.025)	61.9%	38.1%
<i>Worker Rents</i>				
Per-worker Dollars	\$13,469 (788)	\$2,859 (2,662)	70.5%	29.5%
Share of Earnings	17.1% (0.9)			
<i>Firm Rents</i>				
Per-worker Dollars	\$9,403 (1,451)	\$4,763 (3,370)	0.0%	100.0%
Share of Profits	29.2% (2.6)			

Notes. This table presents estimates of wage markdowns and rents. Variance decompositions are calculated via the Law of Total Variance. Standard errors, shown in parentheses, are bootstrapped using 500 bootstrap samples.

VI.B. Misallocation and Inequality

To evaluate the welfare impact of imperfect competition, I compare the equilibrium outcomes in the monopsonistic economy to counterfactual outcomes that would arise in a competitive (Walrasian) economy, where firms do not exert any wage-setting power. In my framework, misallocation occurs because firms vary in their wage-setting power. This variation can lead workers to sort into firms that they would not have chosen in a competitive labor market.

Table 4: Estimates of Counterfactual Comparisons

	Monopsonistic Labor Market	Competitive Labor Market	% Difference in Dollars
Total Welfare (log dollars)	12.71 (0.21)	12.81 (0.23)	9.58%
Mean Earnings (log dollars)	11.29 (0.00)	11.46 (0.01)	15.23%
<u>Variance in Log Earnings</u>			
Total	0.112 (0.00)	0.110 (0.01)	-2.29%
Between Firms	0.028 (0.00)	0.025 (0.01)	-11.12%
Within Firms	0.084 (0.00)	0.085 (0.00)	0.63%

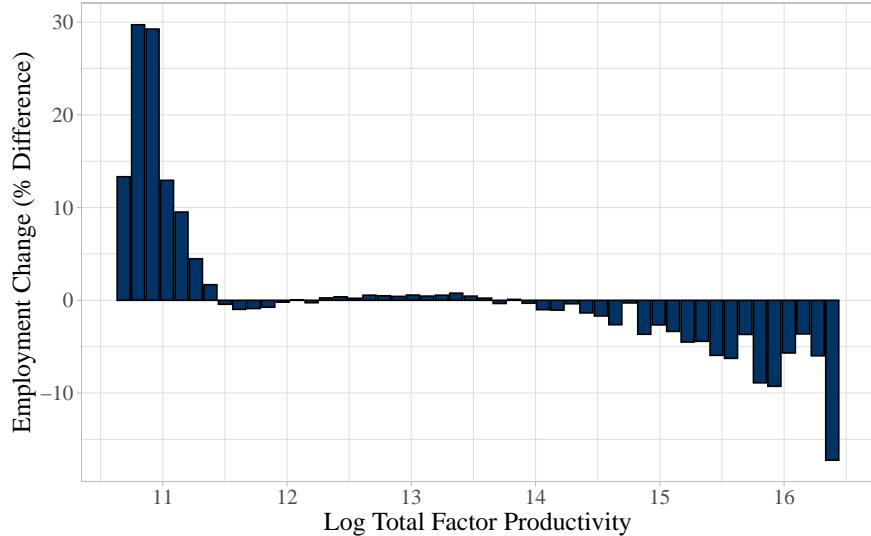
Notes. This table compares outcomes in the monopsonistic economy to counterfactual outcomes in a competitive (Walrasian) economy. The standard errors, shown in parentheses, are bootstrapped using 500 bootstrap samples.

Table 4 presents estimates for counterfactual comparisons. These estimates suggest that imperfect competition produces large misallocation of workers to firms, resulting in a 9.58% reduction in total welfare in dollar terms. Furthermore, in a competitive labor market, average

earnings would increase by 15.23%. However, despite these changes, the overall impact on wage inequality, as measured by the total variance in log earnings, appears to be minimal.

To understand how workers would be reallocated in a competitive labor market, recall that I estimate larger markdowns for firms that have lower effective wages. This suggests that, relative to the competitive benchmark, workers are being allocated away from less productive firms. Indeed, I find that the correlation between a firm's log labor $\ell_{j\tau}$ and its log TFP $t_{j\tau}$ is higher (0.155) in a monopsonistic market than it would be in a competitive market (0.089). Moreover, in Figure 6, I plot the changes in estimated labor shares that would occur if labor wedges were eliminated across the distribution of firm TFP. This figure illustrates that, under the competitive allocation, more workers would be employed at firms with lower productivities. Meanwhile, fewer workers would be employed at firms with higher productivities.

Figure 7: Estimated Reallocation of Labor by Eliminating Labor Wedges



Notes. This figure depicts the reallocation of labor that results from eliminating labor wedges. It plots the estimated percent change in employment across different groups of firms, organized by their estimated log TFPs.

VII. Conclusions

In this paper, I examine how an employer's wage-setting power depends on the characteristics of its workforce, and I assess the consequences of this relationship for welfare and inequality. To this end, I develop a structural model that integrates labor market frictions with individual heterogeneity in workers' trade-offs between wages and non-wage job characteristics. I show that these two factors, taken together, cause employers to face varying labor supply elasticities, resulting in differences in labor market power both within and across firms. Moreover, I show that this framework is empirically tractable and can be used to draw inferences about sorting, imperfect competition, and rent sharing. I estimate the model using a matched worker-firm

panel dataset that covers the universe of workers and firms in Norway between 1995 and 2018.

One key empirical finding from my estimates is that, after accounting for workers' skills, firms with lower wages tend to have more labor market power. My model offers an explanation for this finding. That is, when workers make different trade-offs over job attributes, they tend to sort accordingly, with workers who prioritize amenities (e.g., flexible work hours, fewer physical demands, and/or more societal impact) choosing jobs that pay less. In a labor market that is imperfectly competitive, these sorting patterns lead lower-paying firms to attract labor that is less elastic (i.e., less sensitive to wage changes), thereby increasing wage-setting power. This phenomenon has important welfare implications, because the disparities in market power between firms generate misallocation of labor relative to the competitive benchmark. In a perfectly competitive labor market, more individuals would be employed at low-wage firms.

While this paper offers an initial examination of the relationship between workers' wage-amenity trade-offs and labor market power, there is much room for future work. For example, my framework does not consider the possibility that firms strategically compete for workers by internalizing their market shares. I justify this assumption by testing whether a firm's market share impacts its labor supply elasticities after controlling for wages, finding no significant effect. However, this test does not definitively rule out the existence of strategic interactions. A meaningful extension of my framework would be to examine how a firm's wage-setting power is affected not only by the composition of its workforce but also by its market share.

References

- ABOWD, J. M., F. KRAMARZ, AND D. N. MARGOLIS (1999): “High Wage Workers and High Wage Firms,” *Econometrica*, 67, 251–333.
- ANDREWS, M. J., L. GILL, T. SCHANK, AND R. UPWARD (2008): “High Wage Workers and Low Wage Firms: Negative Assortative Matching or Limited Mobility Bias?” *Journal of the Royal Statistical Society. Series A (Statistics in Society)*, 171, 673–697.
- AZAR, J. A., S. T. BERRY, AND I. MARINESCU (2022): “Estimating Labor Market Power,” Working Paper 30365, National Bureau of Economic Research.
- AZAR, J., AND I. MARINESCU (2024): “Monopsony Power in the Labor Market: From Theory to Policy,” *Annual Review of Economics*, 16, 491–518.
- BERGER, D., K. HERKENHOFF, AND S. MONGEY (2022): “Labor Market Power,” *American Economic Review*, 112, 1147–93.
- BERRY, S., J. LEVINSOHN, AND A. PAKES (1995): “Automobile Prices in Market Equilibrium,” *Econometrica*, 63, 841–890.
- BHASKAR, V., A. MANNING, AND T. TO (2002): “Oligopsony and Monopsonistic Competition in Labor Markets,” *Journal of Economic Perspectives*, 16, 155–174.
- BHULLER, M. S. (2009): “Inndeling av Norge i Arbeidsmarkedsregioner,” *Statistisk Sentralbyrås Notater*, 1–30.
- BOAL, W. M., AND M. R. RANSOM (1997): “Monopsony in the Labor Market,” *Journal of Economic Literature*, 35, 86–112.
- BRUUN, M., J. DELBAS, H. OTTOSSON, AND KEES ZEELENBERG (2021): “Peer Review Report on Compliance with the European Statistics Code of Practice and Further Improvement and Development of the National Statistical System: Norway,” December.
- BURDETT, K., AND D. T. MORTENSEN (1998): “Wage Differentials, Employer Size, and Unemployment,” *International Economic Review*, 39, 257–273.
- CARD, D. (2022): “Who Set Your Wage?” *American Economic Review*, 112, 1075–90.
- CARD, D., A. R. CARDOSO, J. HEINING, AND P. KLINE (2018): “Firms and Labor Market Inequality: Evidence and Some Theory,” *Journal of Labor Economics*, 36, S13–S70.
- CARD, D., J. HEINING, AND P. KLINE (2013): “Workplace Heterogeneity and the Rise of West German Wage Inequality*,” *The Quarterly Journal of Economics*, 128, 967–1015.
- CHAN, M., K. KROFT, E. MATTANA, AND I. MOURIFIÉ (2024): “An Empirical Framework For Matching With Imperfect Competition,” Working Paper 32493, National Bureau of Economic Research.

- CHEN, X. (2007): “Chapter 76 Large Sample Sieve Estimation of Semi-Nonparametric Models,” Volume 6 of *Handbook of Econometrics*: Elsevier, 5549–5632.
- DEB, S., J. EECKHOUT, A. PATEL, AND L. WARREN (2024): “Walras–Bowley Lecture: Market Power and Wage Inequality,” *Econometrica*, 92, 603–636.
- EECKHOUT, J., AND P. KIRCHER (2011): “Identifying Sorting—In Theory,” *The Review of Economic Studies*, 78, 872–906.
- FAGERENG, A., M. MOGSTAD, AND M. RØNNING (2021): “Why Do Wealthy Parents Have Wealthy Children?” *Journal of Political Economy*, 129, 703–756.
- H. DE FRAHAN, L., T. LAMADON, AND T. MELING (2024): “Why Do Larger Firms Have Lower Labor Shares?”,
- GIBBONS, R., AND L. KATZ (1992): “Does Unmeasured Ability Explain Inter-Industry Wage Differentials?,” *The Review of Economic Studies*, 59, 515–535.
- GIBBONS, R., L. F. KATZ, T. LEMIEUX, AND D. PARENT (2005): “Comparative Advantage, Learning, and Sectoral Wage Determination,” *Journal of Labor Economics*, 23, 681–724.
- HAANWINCKEL, D. (2023): “Supply, Demand, Institutions, and Firms: A Theory of Labor Market Sorting and the Wage Distribution,” Working Paper 31318, National Bureau of Economic Research.
- HALL, R. E., AND A. I. MUELLER (2018): “Wage Dispersion and Search Behavior: The Importance of Nonwage Job Values,” *Journal of Political Economy*, 126, pp. 1594–1637.
- HAMERMESH, D. S. (1999): “Changing Inequality in Markets for Workplace Amenities,” *The Quarterly Journal of Economics*, 114, 1085–1123.
- JAROSCH, G., J. S. NIMCZIK, AND I. SORKIN (2024): “Granular Search, Market Structure, and Wages,” *The Review of Economic Studies*, rdae004.
- KLINE, P. M. (2025): “Labor Market Monopsony: Fundamentals and Frontiers,” Working Paper 33467, National Bureau of Economic Research.
- KROFT, K., Y. LUO, M. MOGSTAD, AND B. SETZLER (2025): “Imperfect Competition and Rents in Labor and Product Markets: The Case of the Construction Industry,” *American Economic Review*, 115, 2926–69, 10.1257/aer.20220577.
- LAMADON, T., M. MOGSTAD, AND B. SETZLER (2022): “Imperfect Competition, Compensating Differentials, and Rent Sharing in the US Labor Market,” *American Economic Review*, 112, 169–212.
- LISE, J., AND F. POSTEL-VINAY (2020): “Multidimensional Skills, Sorting, and Human Capital Accumulation,” *American Economic Review*, 110, 2328–76.

- MAESTAS, N., K. J. MULLEN, D. POWELL, T. VON WACHTER, AND J. B. WENGER (2017): “Working Conditions in the United States: Results of the 2015 American Working Conditions Survey,” Technical report, Santa Monica, CA: RAND Corporation.
- (2023): “The Value of Working Conditions in the United States and the Implications for the Structure of Wages,” *American Economic Review*, 113, 2007–47.
- MANNING, A. (2010): “The plant size-place effect: agglomeration and monopsony in labour markets,” *Journal of Economic Geography*, 10, 717–744.
- (2011): “Chapter 11 - Imperfect Competition in the Labor Market,” Volume 4 of Handbook of Labor Economics: Elsevier, 973–1041.
- (2021): “Monopsony in Labor Markets: A Review,” *ILR Review*, 74, 3–26.
- MAS, A., AND A. PALLAIS (2017): “Valuing Alternative Work Arrangements,” *American Economic Review*, 107, 3722–59.
- MURPHY, K. M., AND R. H. TOPEL (1990): *Efficiency Wages Reconsidered: Theory and Evidence*, 204–240, London: Palgrave Macmillan UK.
- NEVO, A. (2000): “A Practitioner’s Guide to Estimation of Random-Coefficients Logit Models of Demand,” *Journal of Economics & Management Strategy*, 9, 513–548.
- PIERCE, B. (2001): “Compensation Inequality*,” *The Quarterly Journal of Economics*, 116, 1493–1525.
- ROBINSON, J. (1933): *Journal of the Royal Statistical Society*, 97, 671–674.
- ROSEN, S. (1986): “Chapter 12 The theory of equalizing differences,” Volume 1 of Handbook of Labor Economics: Elsevier, 641–692.
- ROUSSILLE, N., AND B. SCUDERI (2023): “Bidding for Talent: A Test of Conduct in a High-Wage Labor Market,” *iza discussion papers*, Institute of Labor Economics (IZA).
- SHARMA, G. (2023): “Monopsony and Gender.”
- SOKOLOVA, A., AND T. SORENSEN (2021): “Monopsony in Labor Markets: A Meta-Analysis,” *ILR Review*, 74, 27–55.
- SONG, J., D. J. PRICE, F. GUVENEN, N. BLOOM, AND T. VON WACHTER (2018): “Firming Up Inequality,” *The Quarterly Journal of Economics*, 134, 1–50.
- SORKIN, I. (2018): “Ranking Firms Using Revealed Preference*,” *The Quarterly Journal of Economics*, 133, 1331–1393.
- TABER, C., AND R. VEJLIN (2020): “Estimation of a Roy/Search/Compensating Differential Model of the Labor Market,” *Econometrica*, 88, pp. 1031–1069.

VERBOVEN, F. (1996): “The Nested Logit Model and Representative Consumer Theory,” *Economics Letters*, 50, 57–63.

WIPER, M., D. R. INSUA, AND F. RUGGERI (2001): “Mixtures of Gamma Distributions with Applications,” *Journal of Computational and Graphical Statistics*, 10, 440–454.

WISWALL, M., AND B. ZAFAR (2018): “Preference for the Workplace, Investment in Human Capital, and Gender,” *The Quarterly Journal of Economics*, 133, 457–507.